MIND: In-Network Memory Management for Disaggregated Data Centers

Seung-seob Lee, Yanpeng Yu, Yupeng Tang,
Anurag Khandelwal, Lin Zhong, Abhishek Bhattacharjee

Yale University

Resource Disaggregation

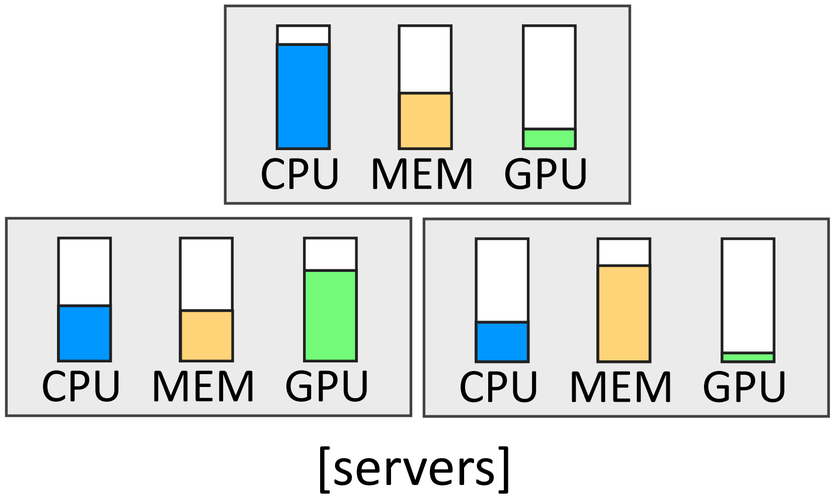

Resource Disaggregation

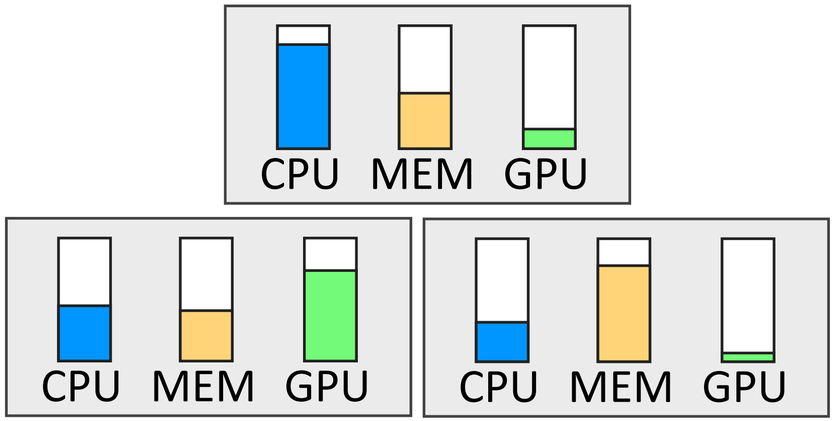

[servers]

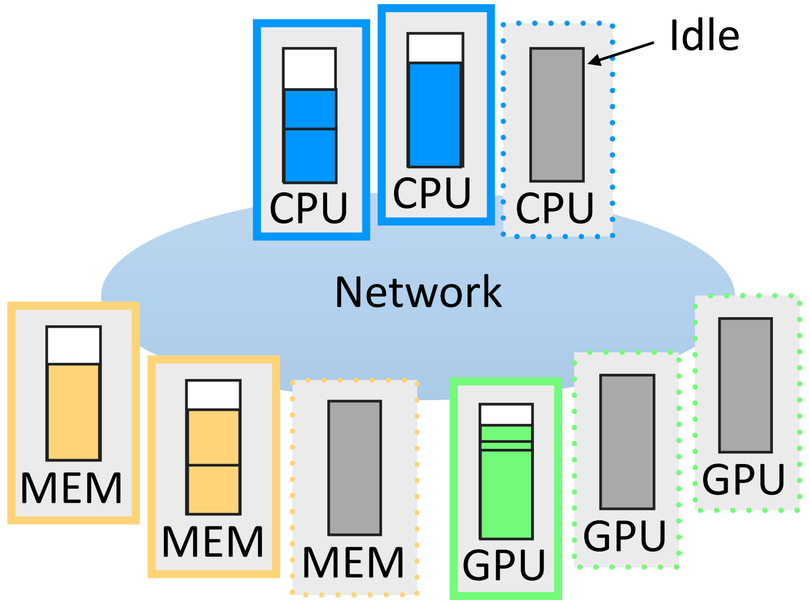

[disaggregated blades]

Resource Disaggregation

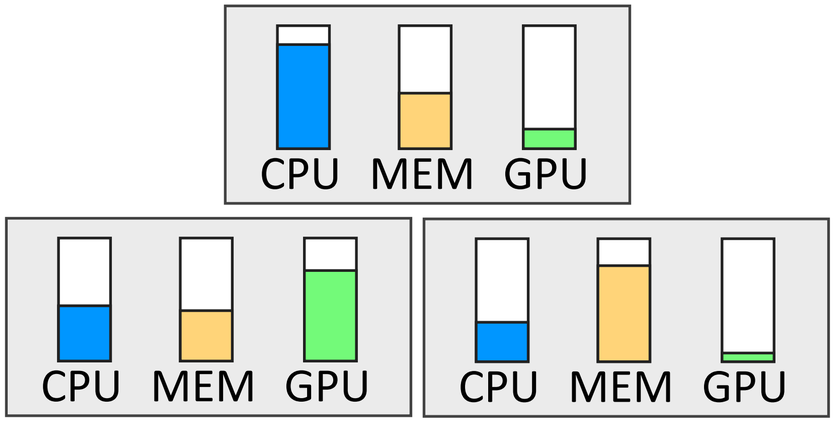

[servers]

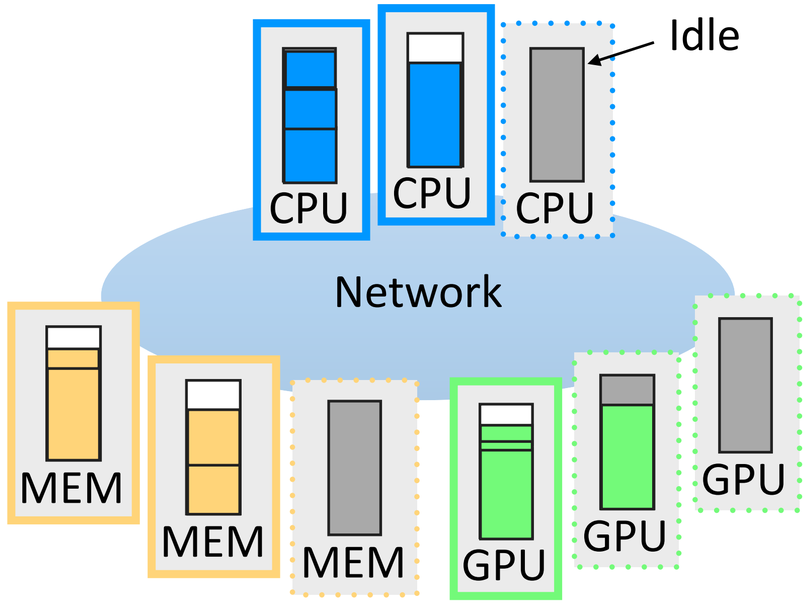

[disaggregated blades]

- **Benefits of resource disaggregation**
  - High resource utilization
  - Easy to manage (modularity)
  - Elastic scalability

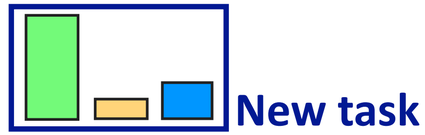

Memory Disaggregation

- **Need for memory disaggregation**

Low utilization
(as low as 30%)

Unbalanced usage
(> 70% of the time in clusters)

Energy consumption
(up to 46% of avg. system)

Memory Disaggregation

- **Need for memory disaggregation**

Low utilization
(as low as 30%)

Unbalanced usage
(> 70% of the time in clusters)

Energy consumption
(up to 46% of avg. system)

- **Timely problem**

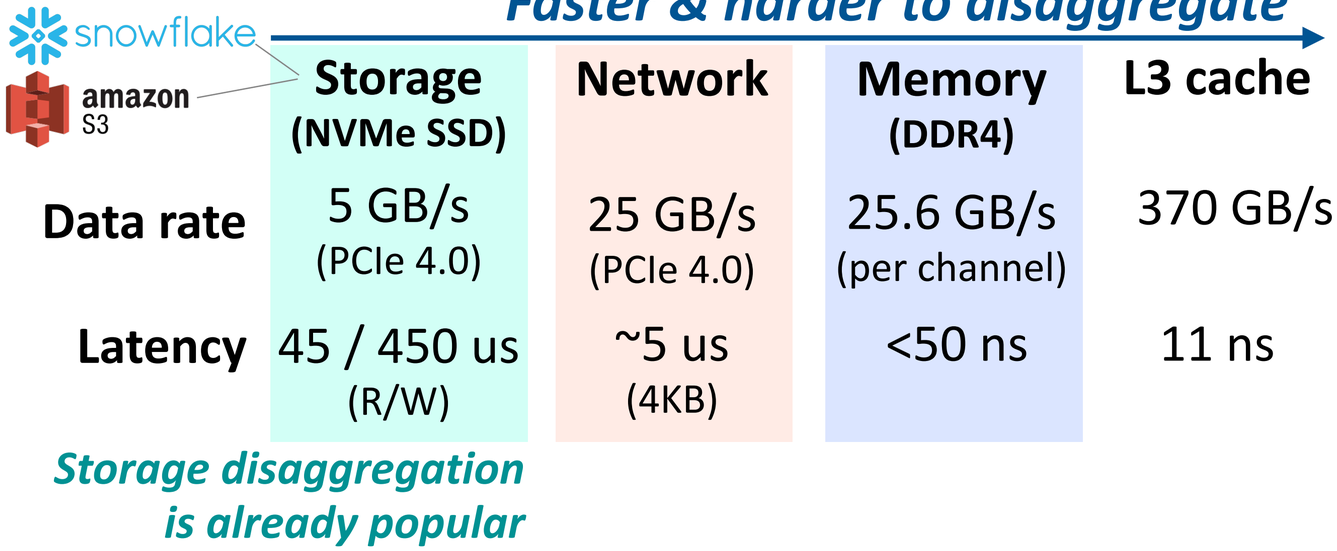

Memory Disaggregation Goals

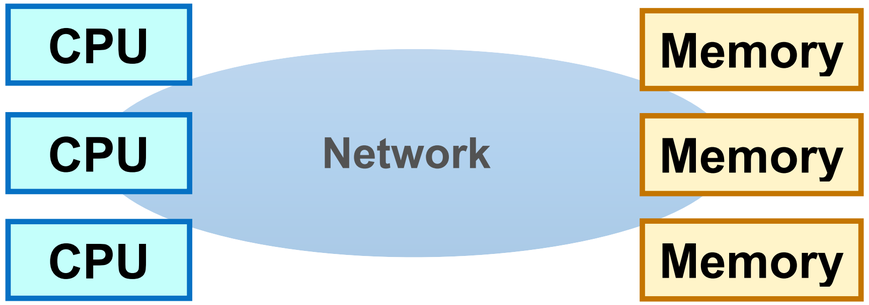

[disaggregated blades]

Memory Disaggregation Goals

- **Performance (CPU ↔ memory)**: high throughput & low latency

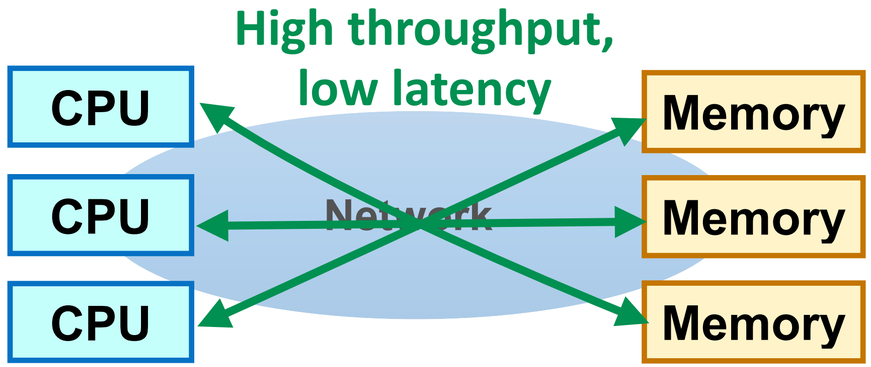

[disaggregated blades]

Memory Disaggregation Goals

- **Performance (CPU ↔ memory)**: high throughput & low latency
- **Transparent elasticity**: flexible resource allocation

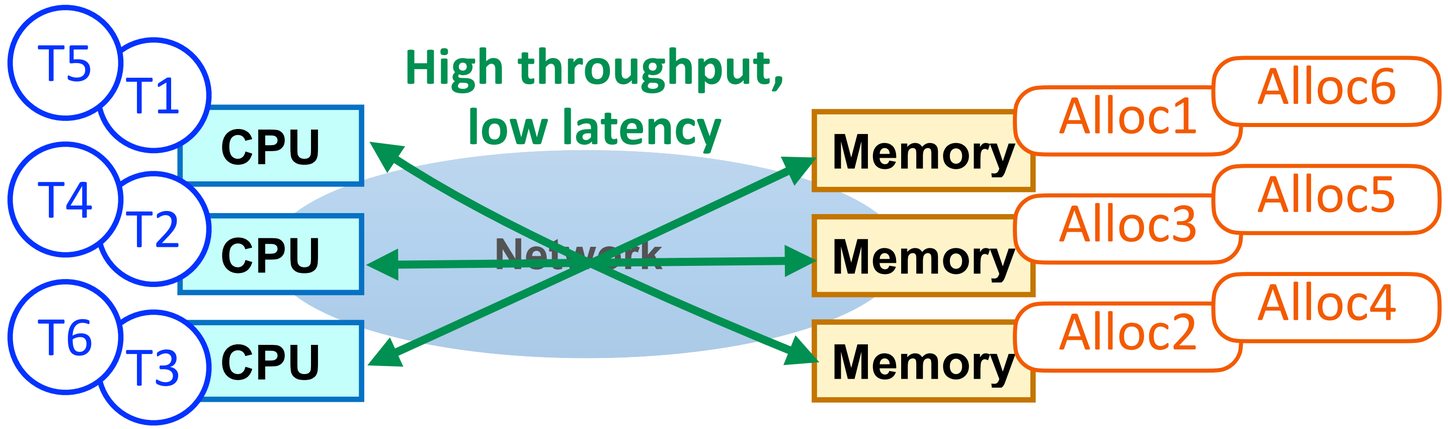

[disaggregated blades]

Existing Approaches

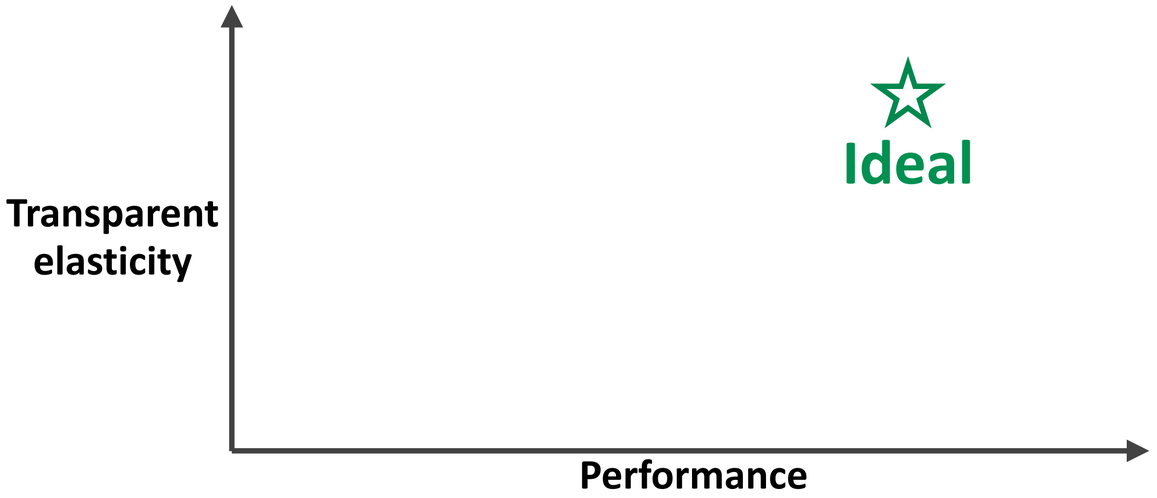

Existing Approaches

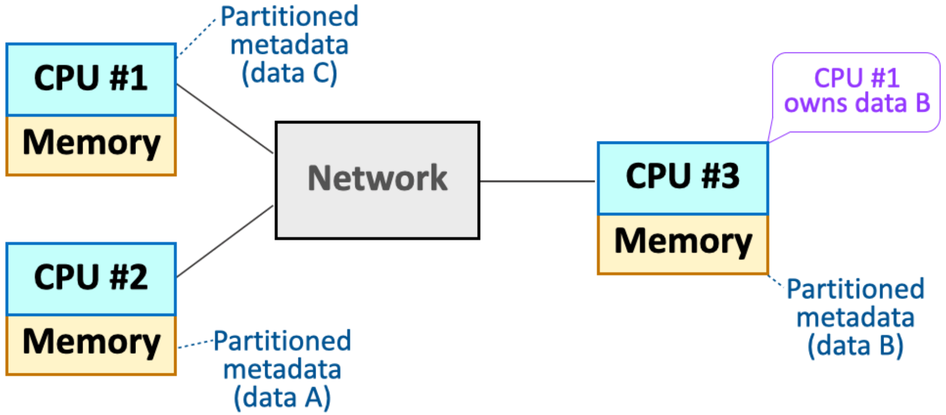

Distributed shared memory (DSM)

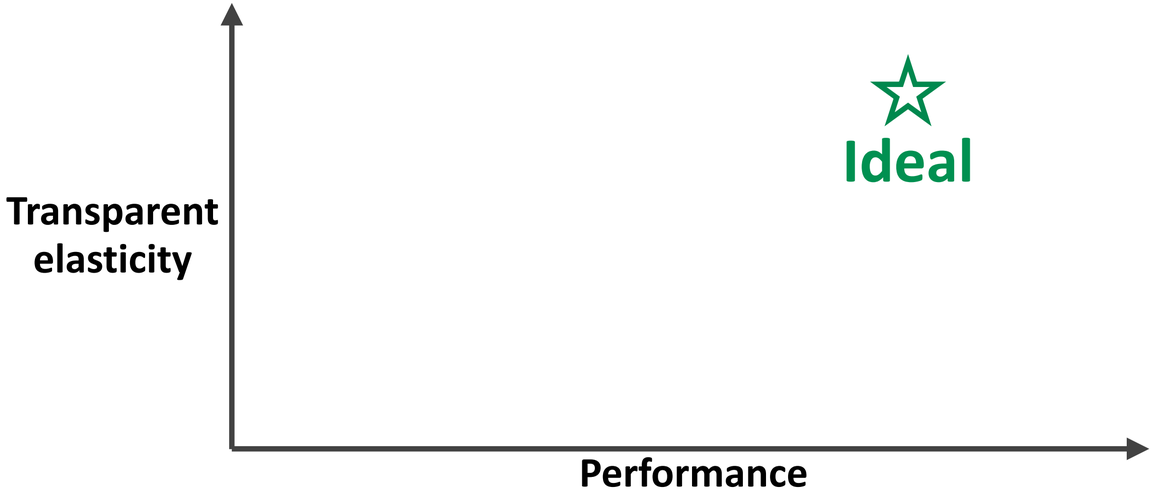

Existing Approaches

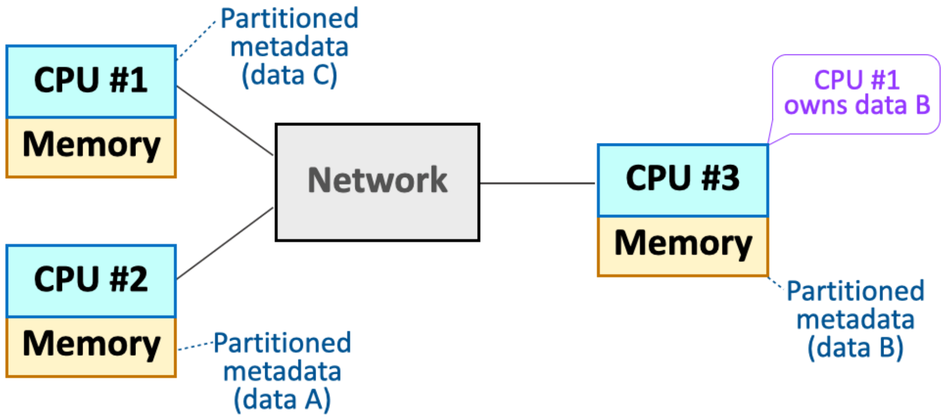

Distributed shared memory (DSM)

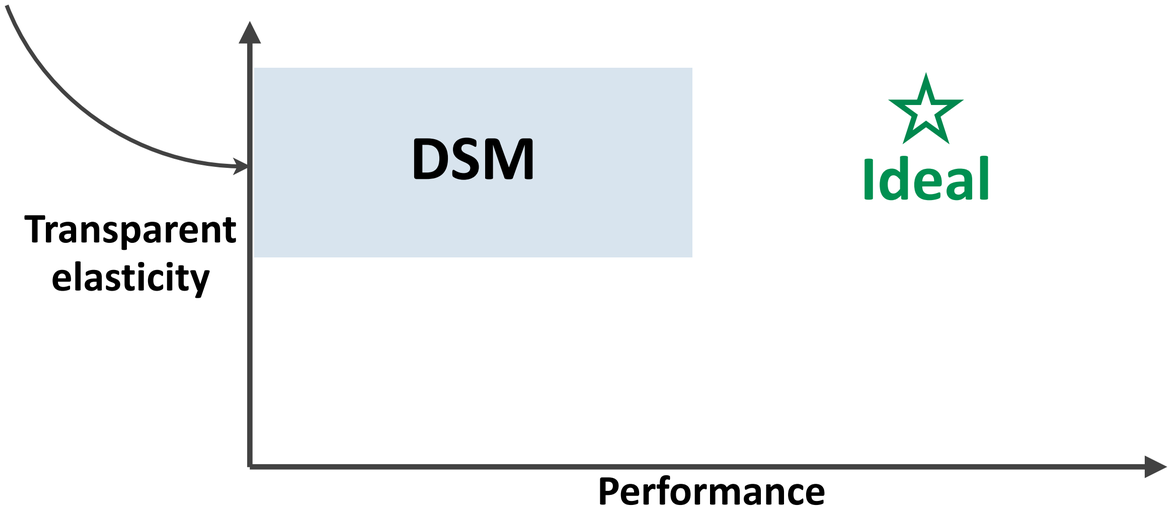

Existing Approaches

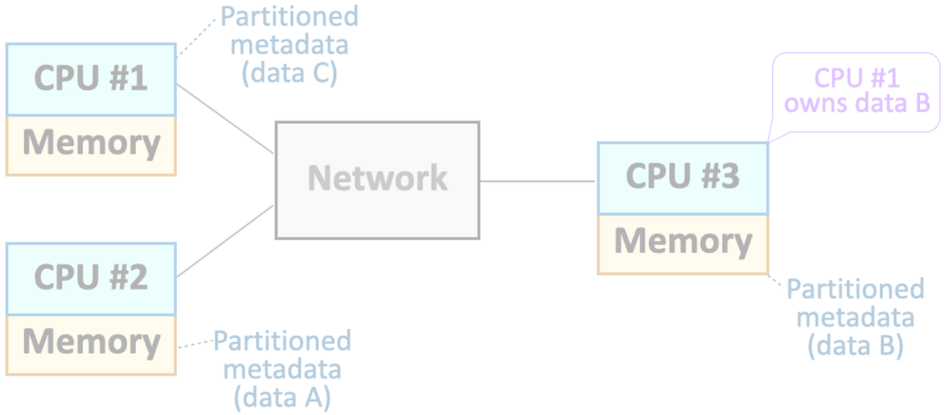

Distributed shared memory (DSM)

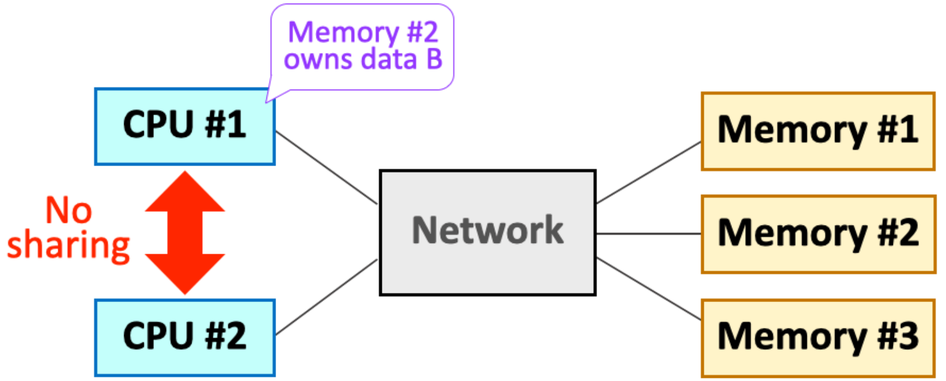

Recent disaggregated memory schemes

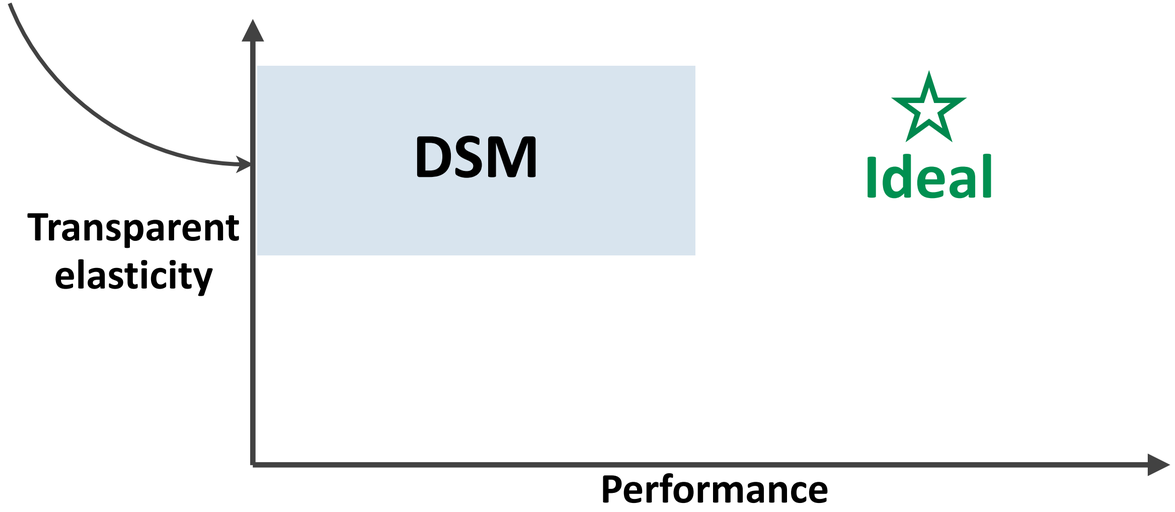

Existing Approaches

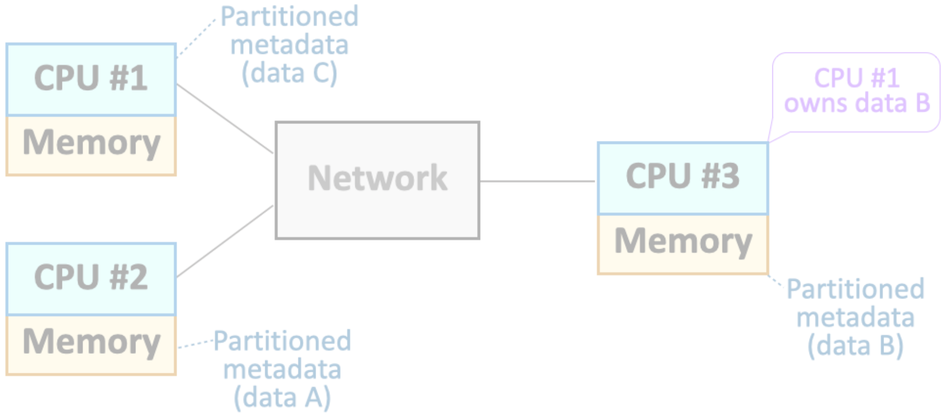

Distributed shared memory (DSM)

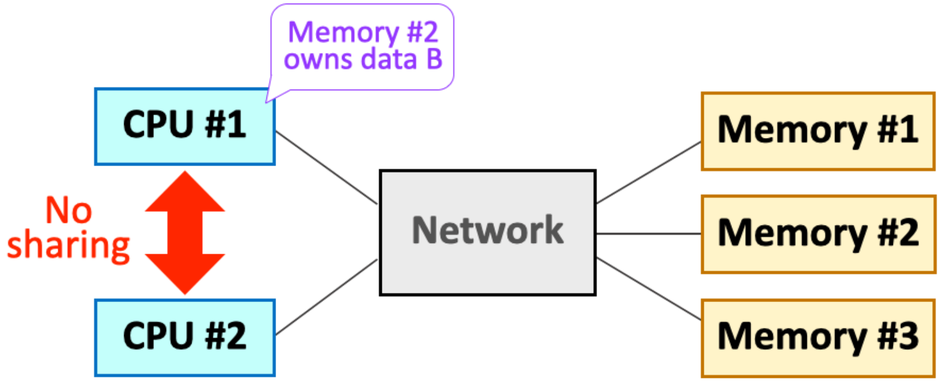

Recent disaggregated memory schemes

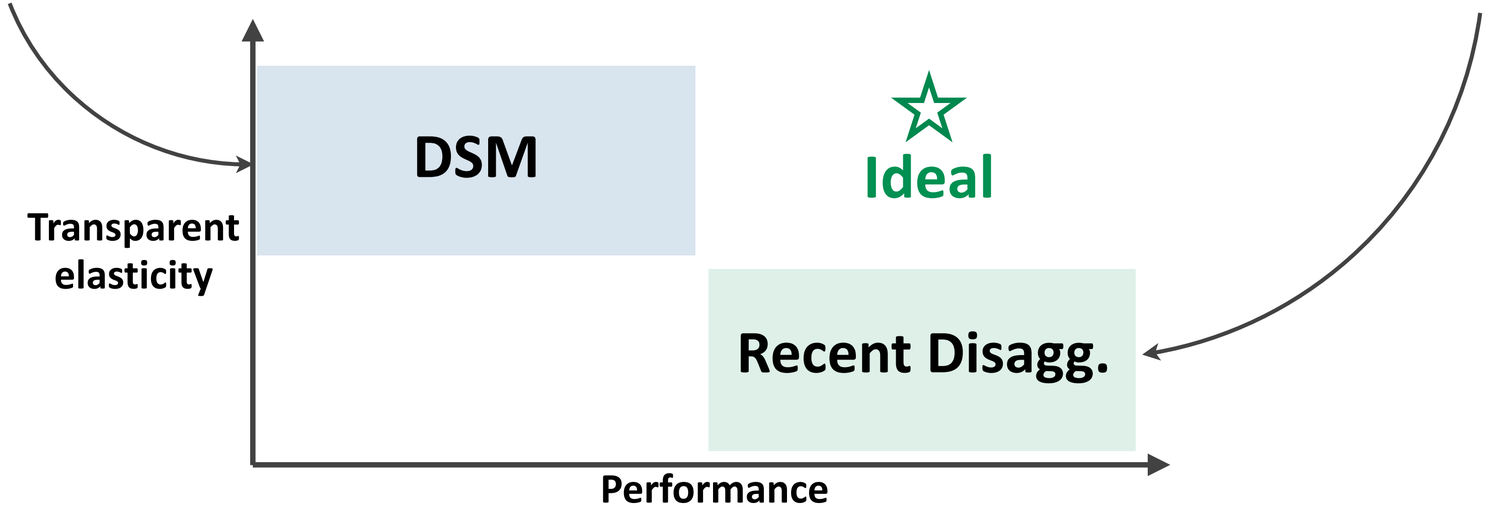

Key Insight — In-network Memory management

- **Central location:** global view & processing directly in the data path

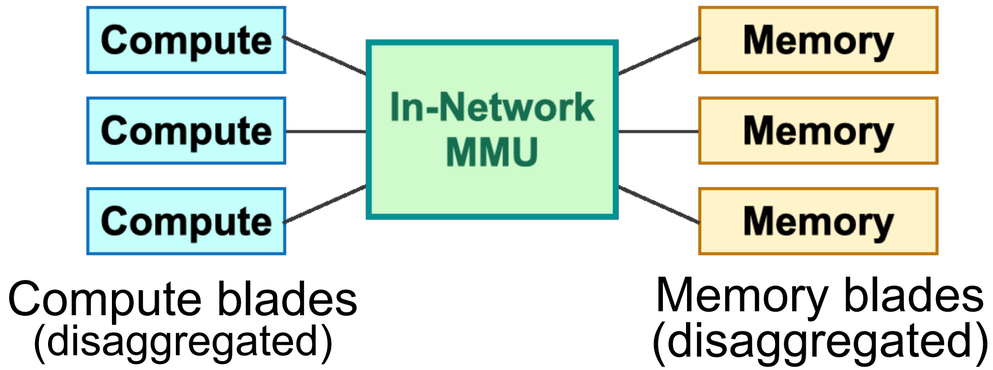

Key Insight — In-network Memory management

- **Central location:** global view & processing directly in the data path
- **Programmable switching ASIC:** flexible processing at line rate

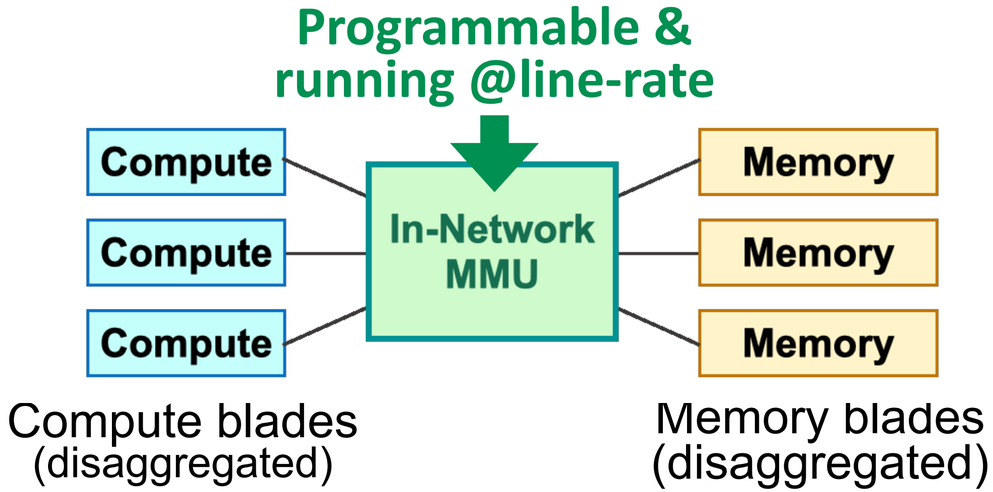

Key Insight — In-network Memory management

- **Central location:** global view & processing directly in the data path
- **Programmable switching ASIC:** flexible processing at line rate
- **Similarity between network functions and memory management**

Networking		Memory management
• IP forwarding • IP assignment • Access control • Multicast/broadcast	⟷	• Address translation • Memory allocation • Memory protection • Cache invalidations

MIND: In-network Memory Management

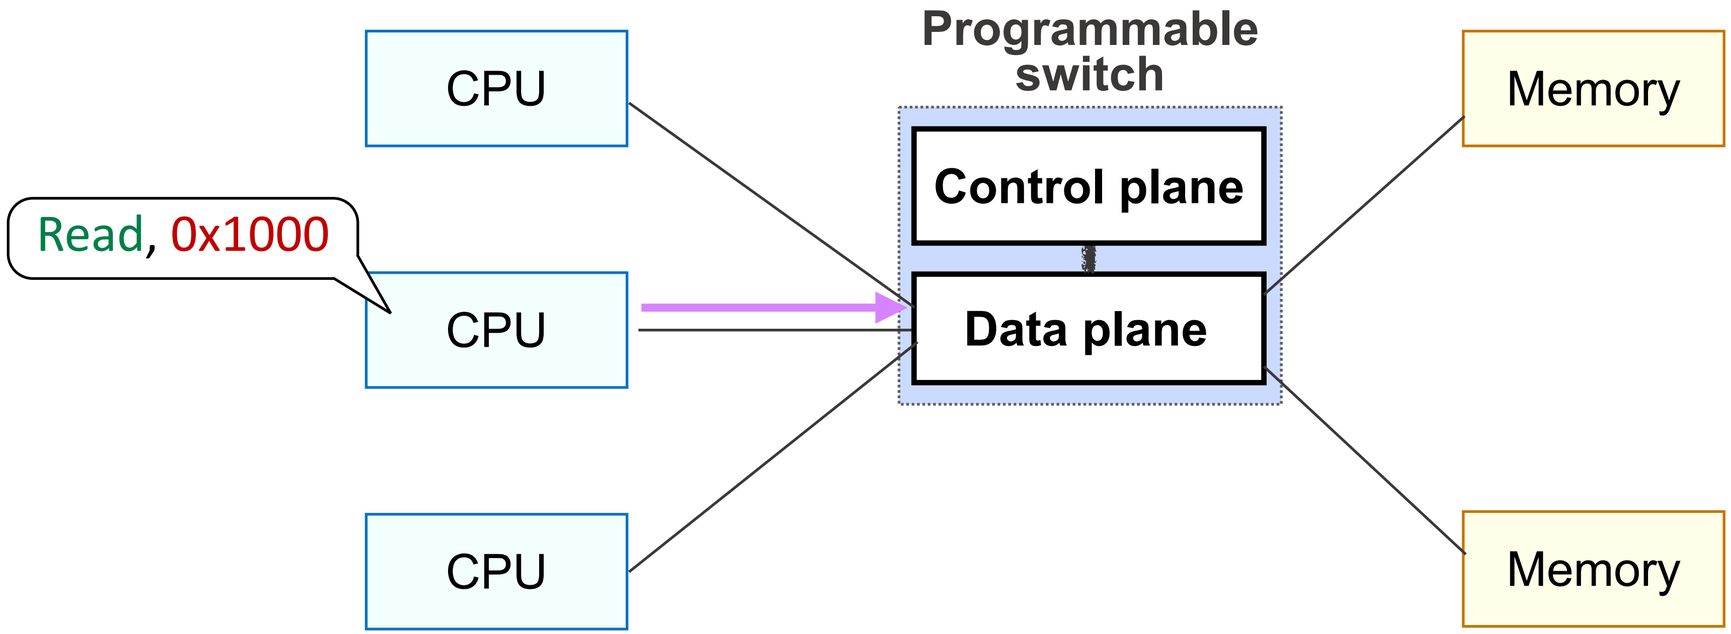

MIND: In-network Memory Management

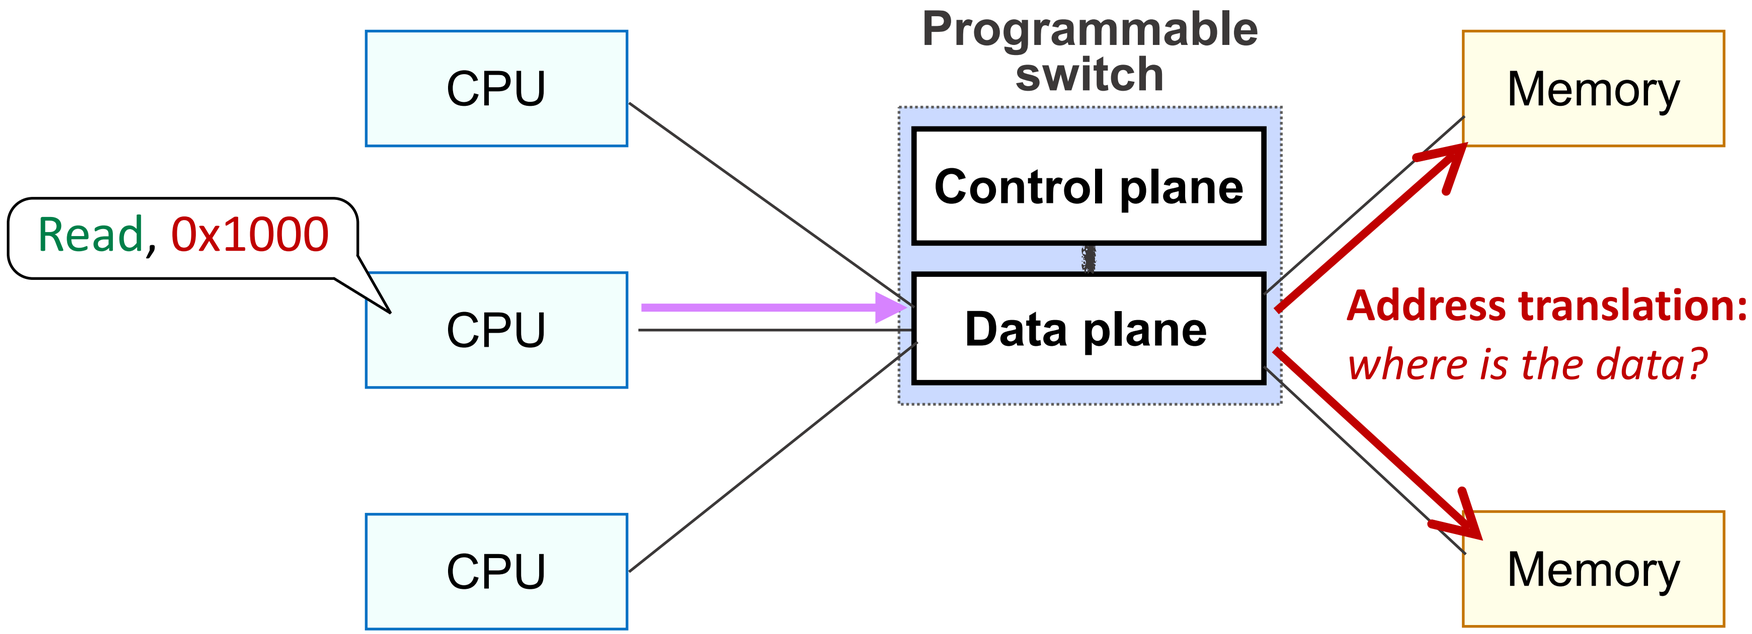

MIND: In-network Memory Management

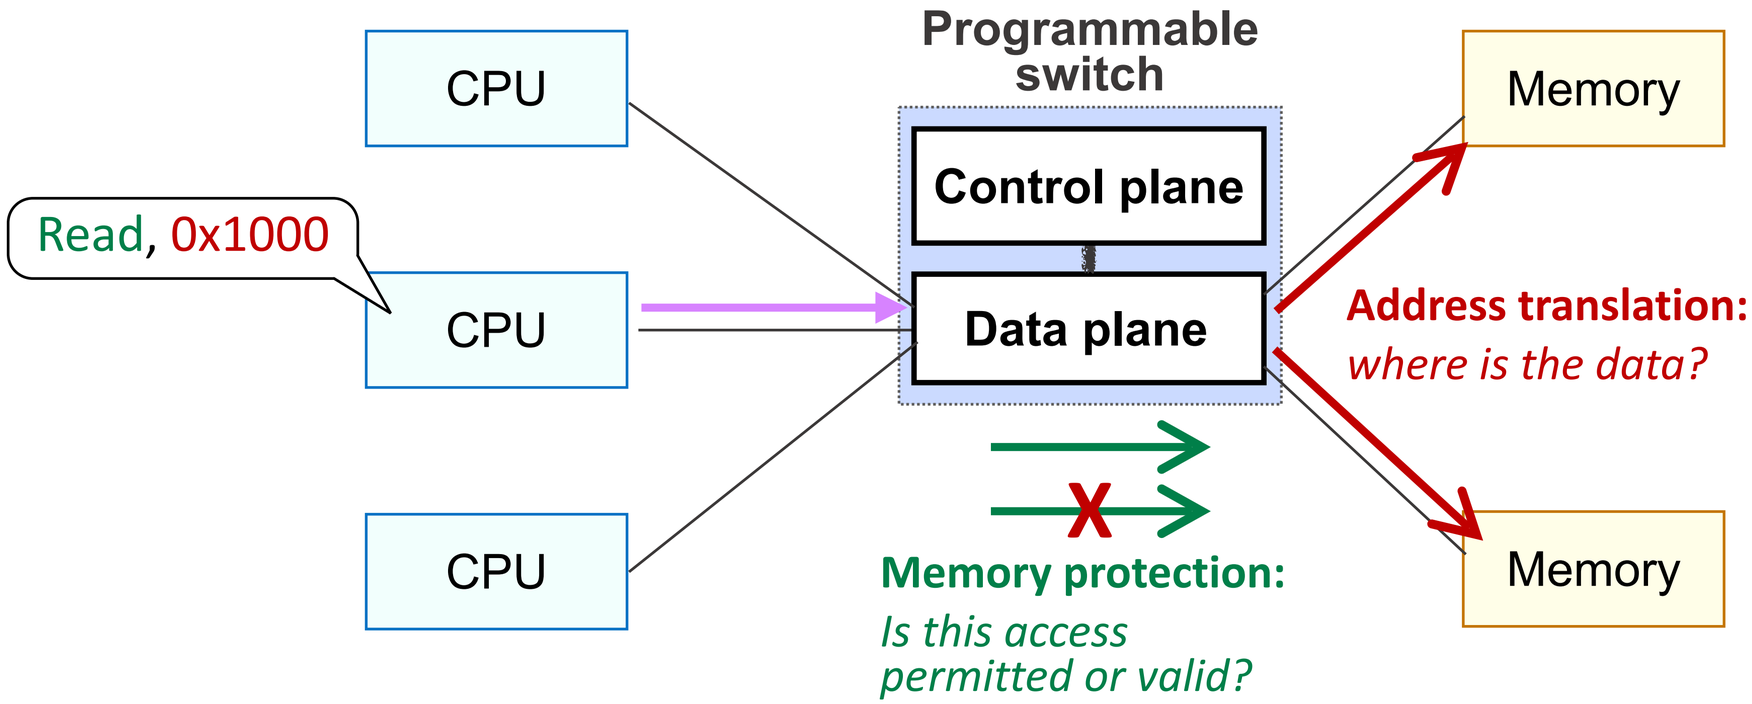

MIND: In-network Memory Management

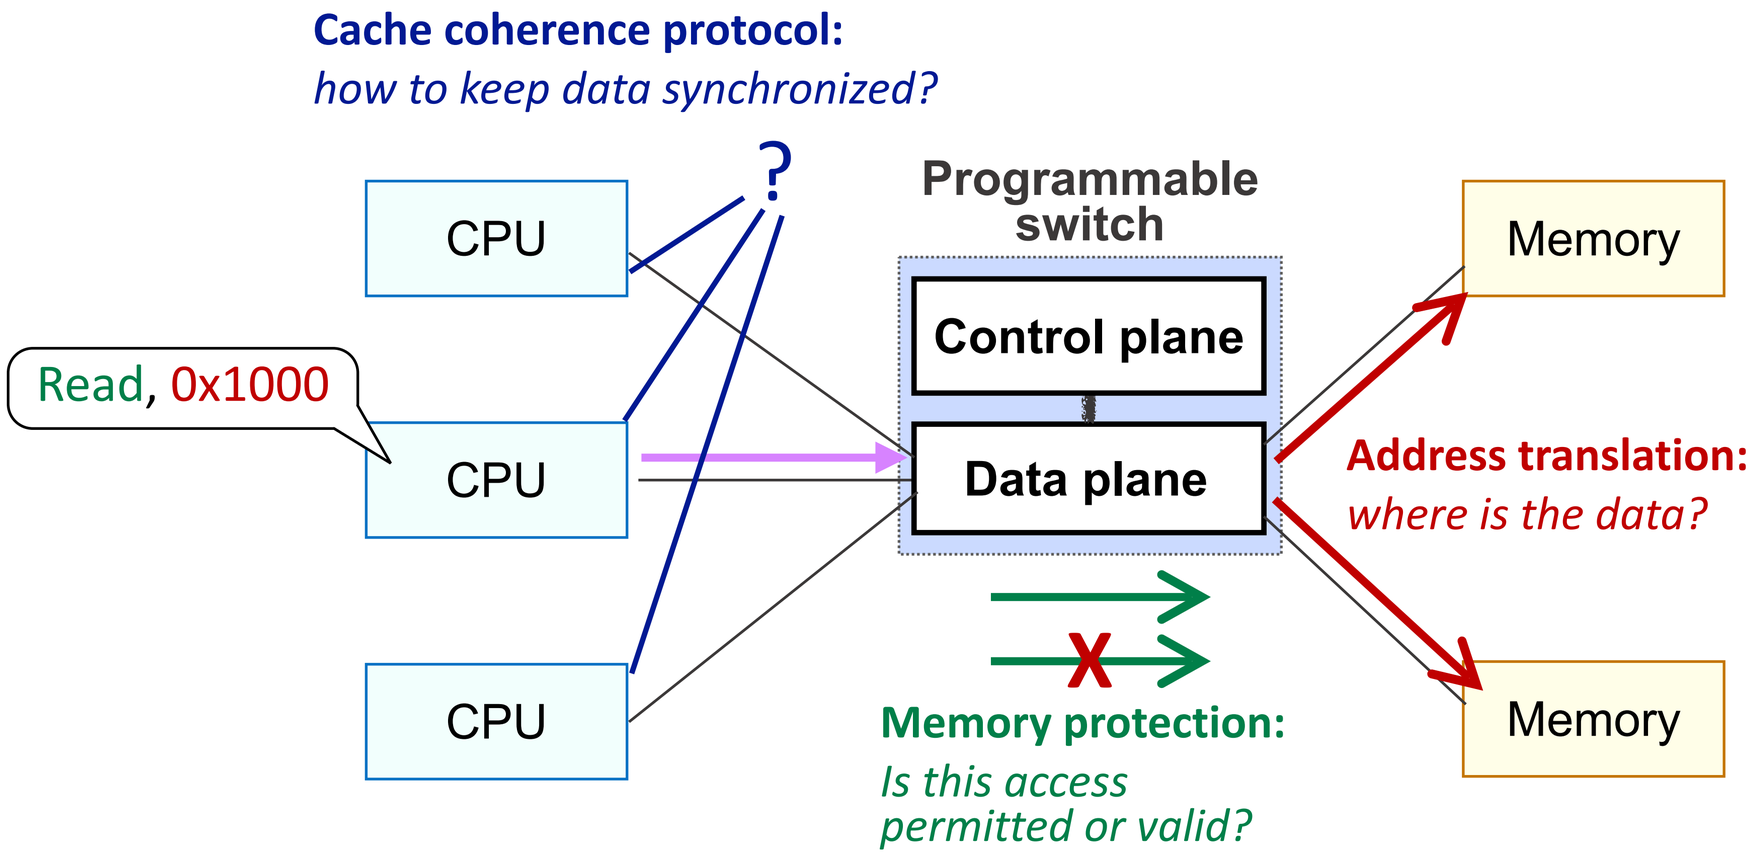

MIND: In-network Memory Management

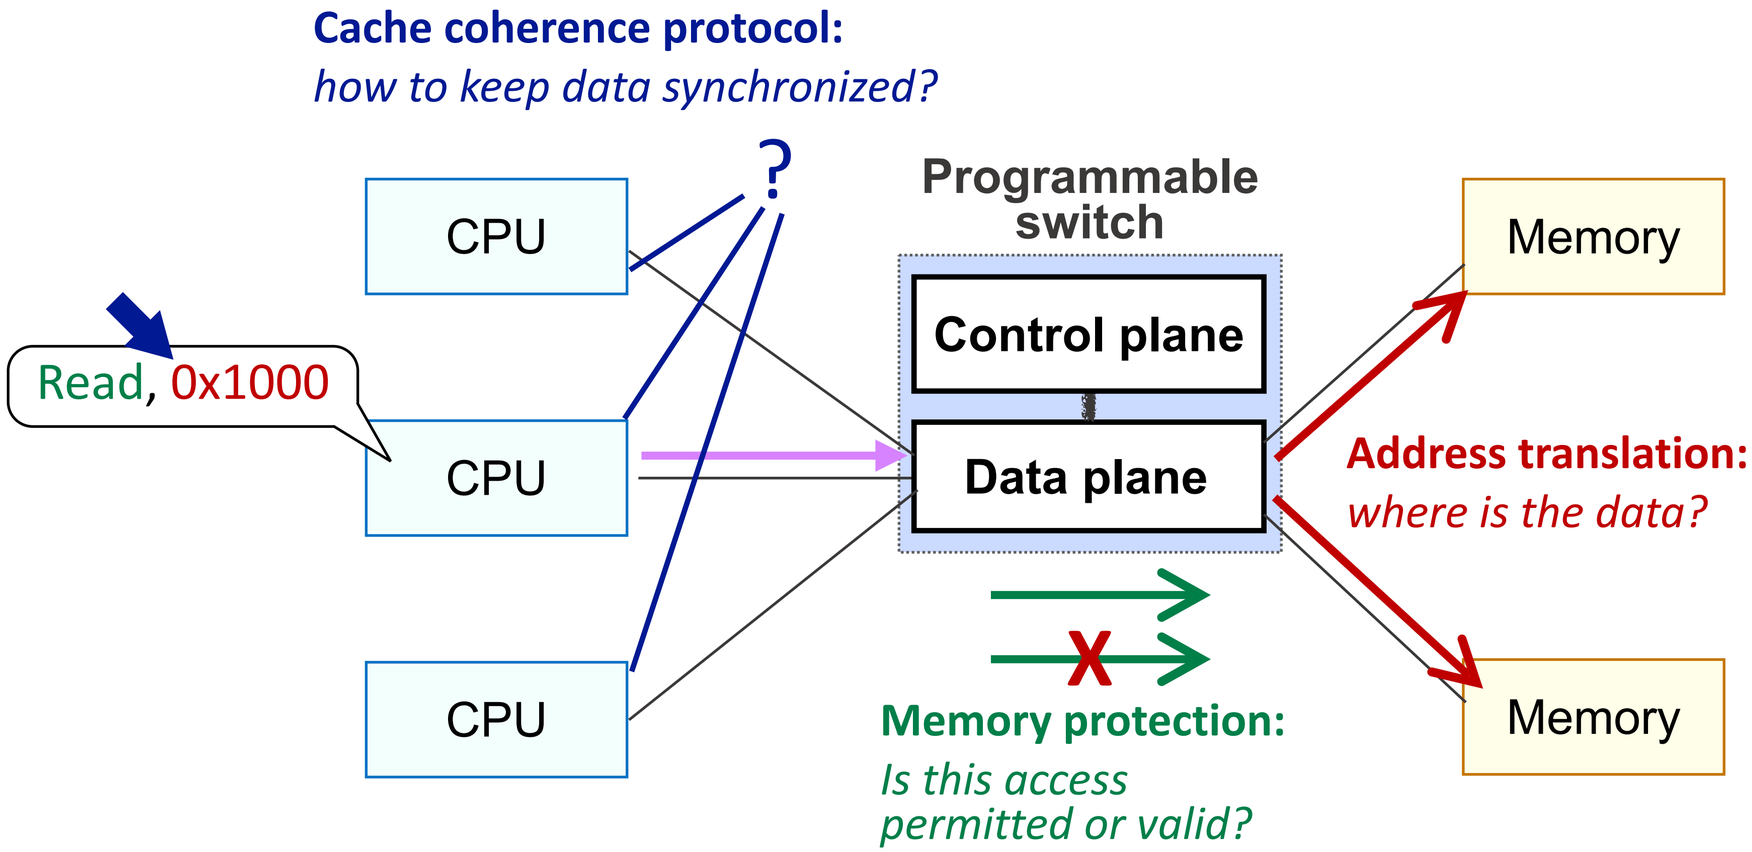

Challenges in Building In-network MMU

- **Limited amount of in-network resources**
  - *Limited size of in-network memory*
    - Tens of MB → not sufficient to store metadata in a traditional way
      - *E.g., page table: 4 KB for 4GB of memory → 1M entries*

Challenges in Building In-network MMU

- **Limited amount of in-network resources**
  - *Limited size of in-network memory*
    - Tens of MB → not sufficient to store metadata in a traditional way
      - *E.g., page table: 4 KB for 4GB of memory → 1M entries*
  - *Limited computation capability*
    - Switching ASIC → not sufficient to directly port traditional MMU functions
      - *E.g., limited number of operations for line rate processing*
        ↔ complicated cache coherence protocol

3 Principles for System Design

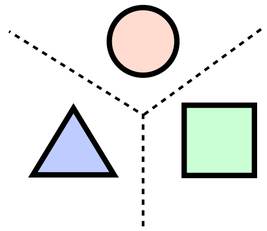

- **P1 - Decouple memory management functionalities**
  → *Each function has the data structure suitable to its purpose*

3 Principles for System Design

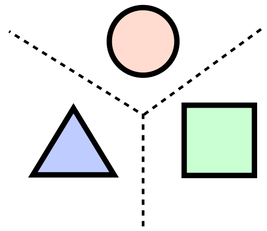

- **P1 - Decouple memory management functionalities**
  → *Each function has the data structure suitable to its purpose*

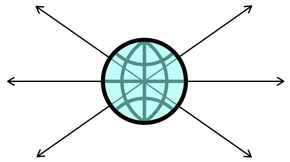

- **P2 - Leverage global view of network**
  → *Make better decisions for memory management*

3 Principles for System Design

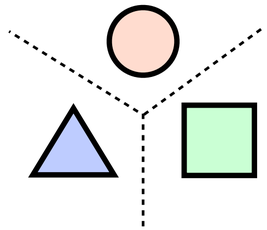

- **P1 - Decouple memory management functionalities**
  → *Each function has the data structure suitable to its purpose*

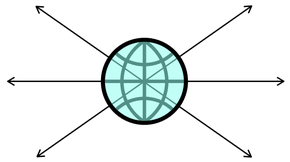

- **P2 - Leverage global view of network**
  → *Make better decisions for memory management*

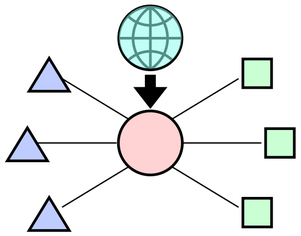

- **P3 - Exploit network-centric hardware primitives**
  → *Reuse network hardware highly optimized for network functions*

MIND overview: challenges and solutions

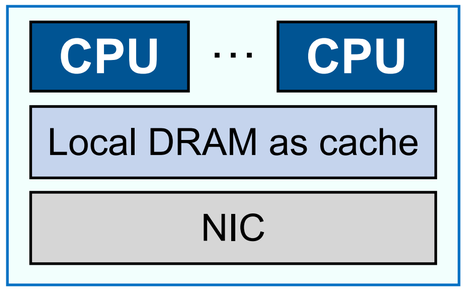

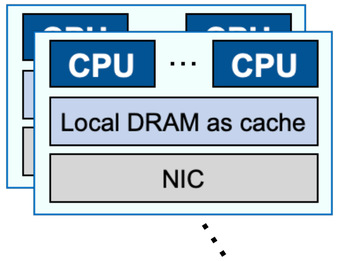

Compute blades

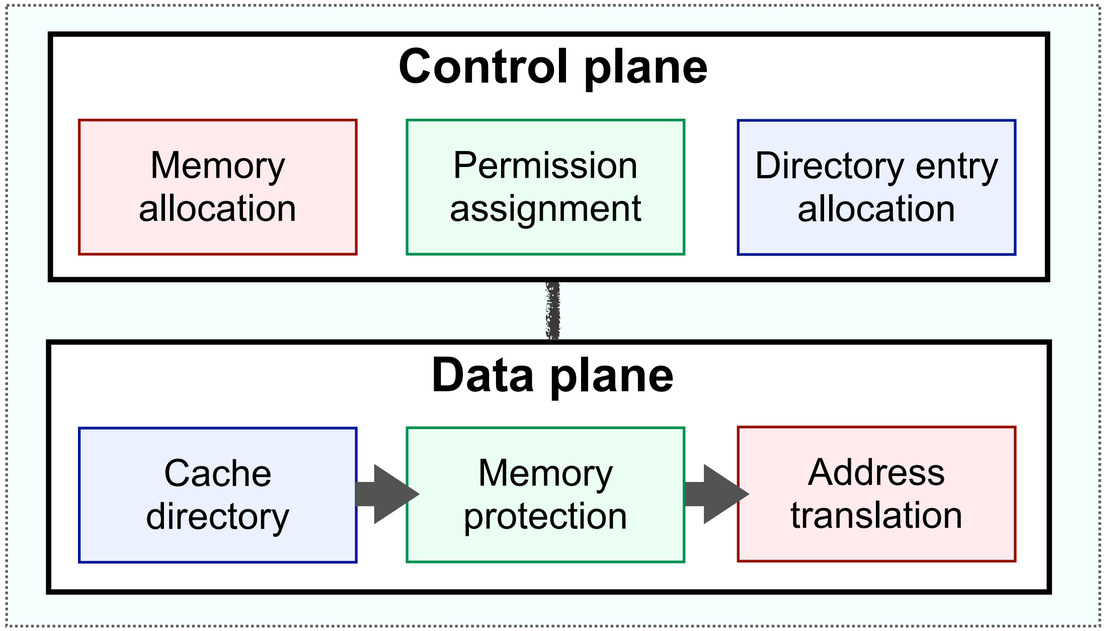

Programmable switch

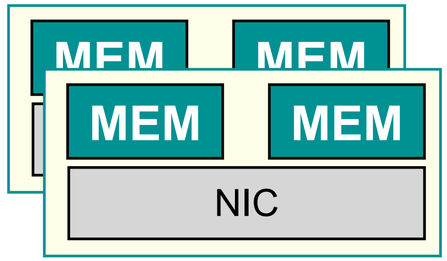

Memory blades

MIND overview: challenges and solutions

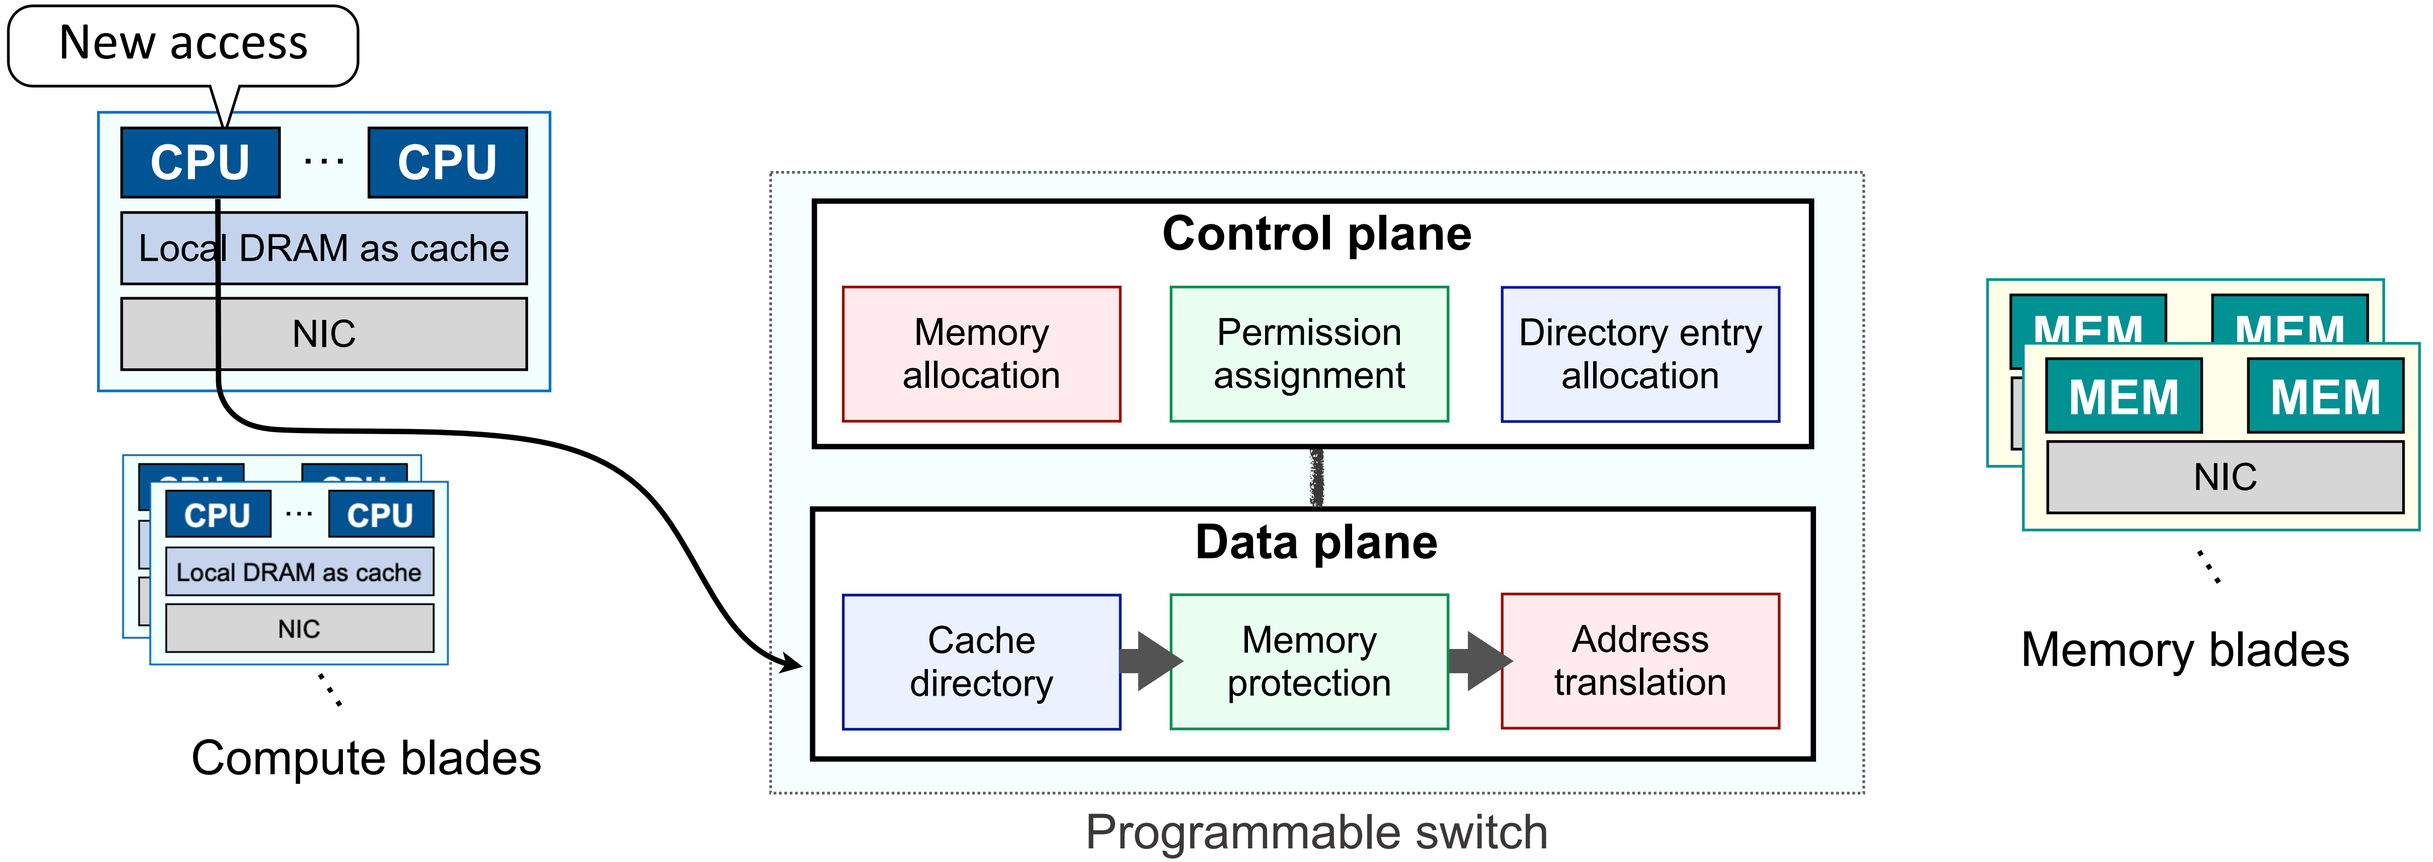

MIND overview: challenges and solutions

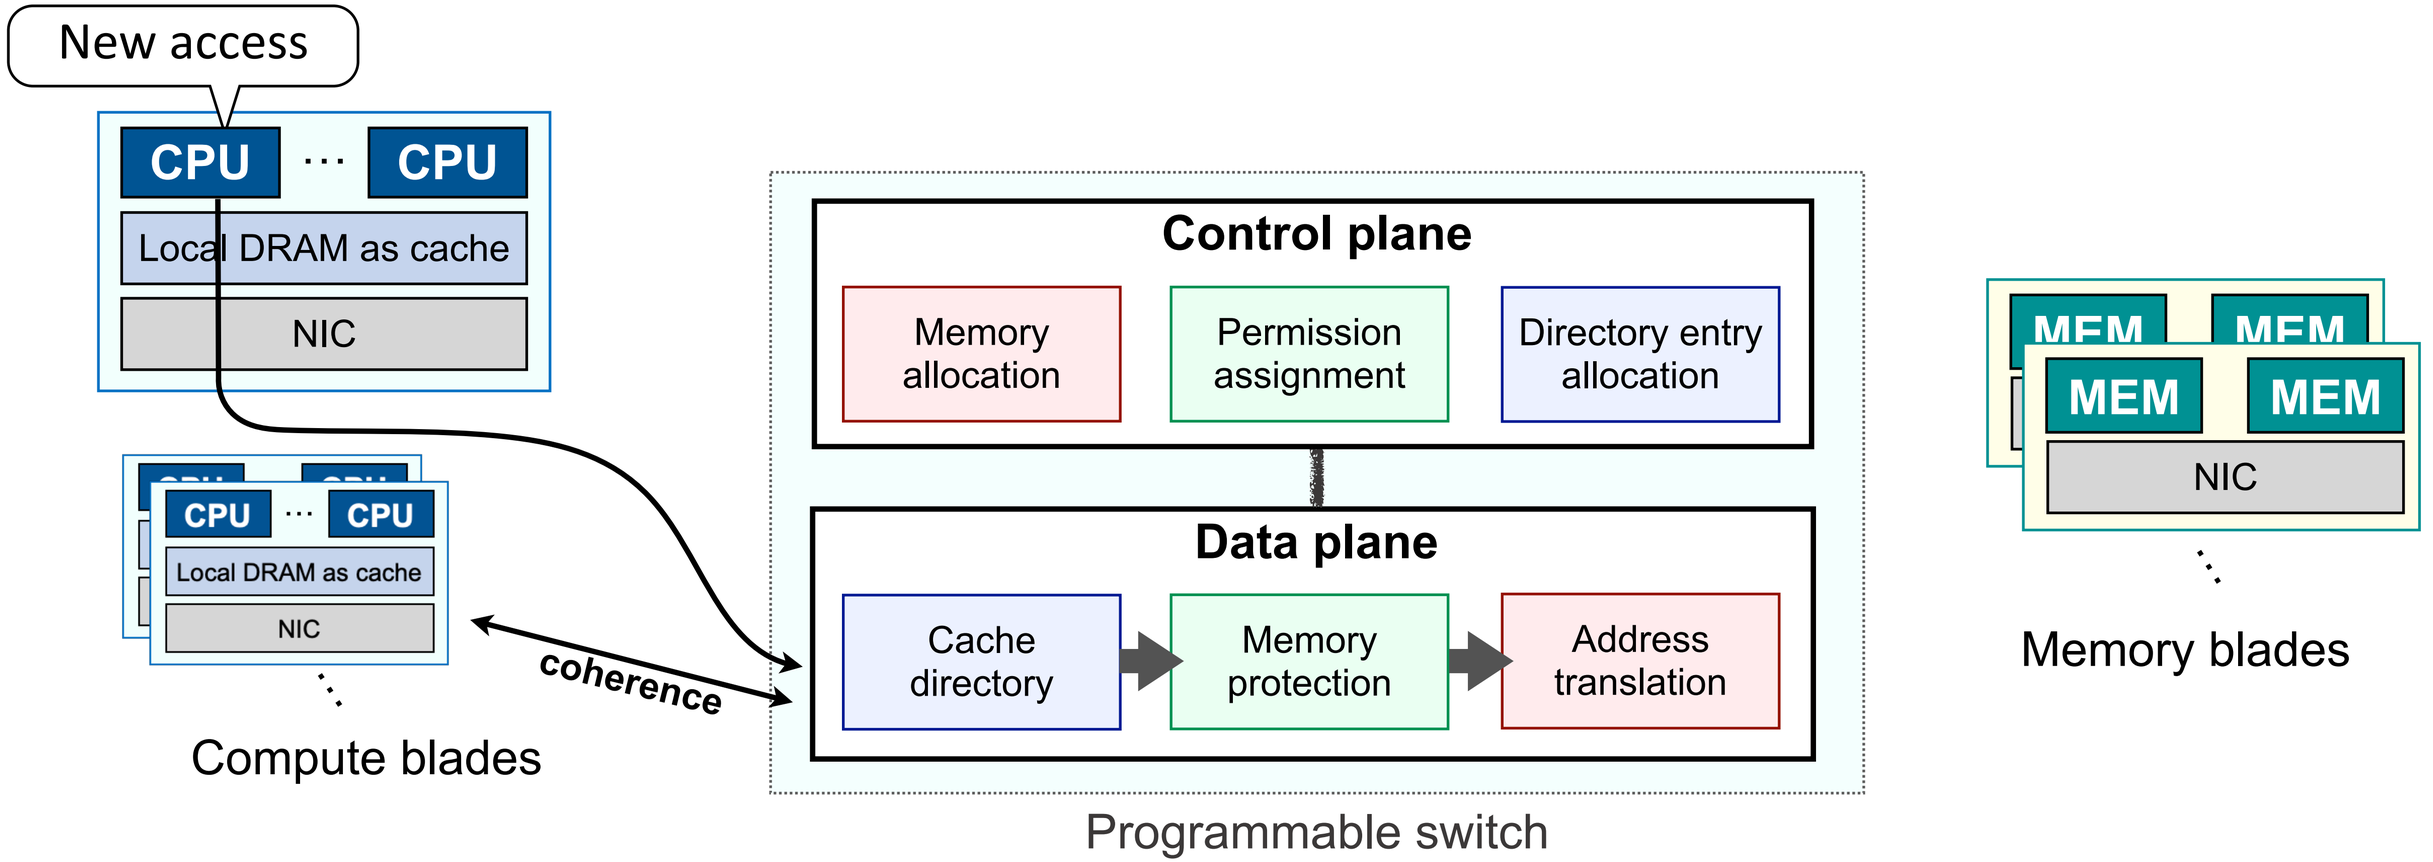

MIND overview: challenges and solutions

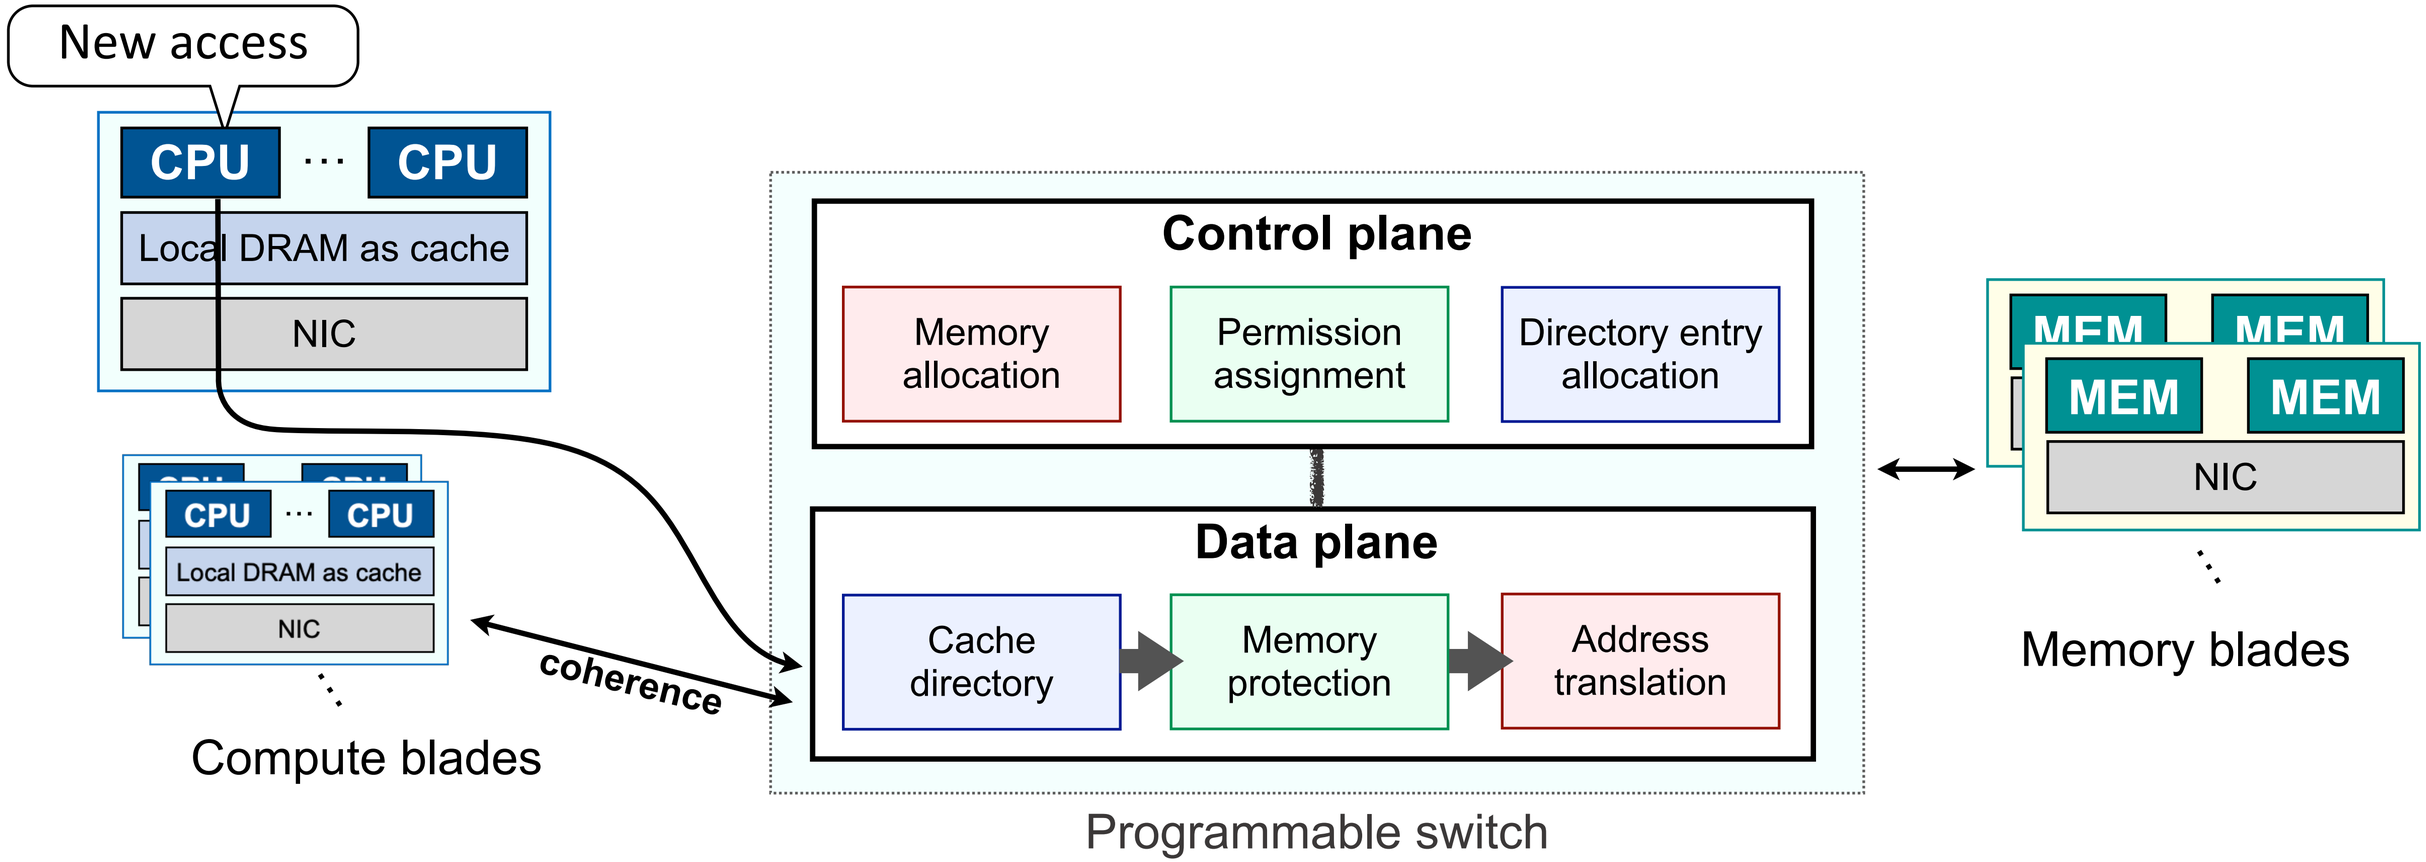

MIND overview: challenges and solutions

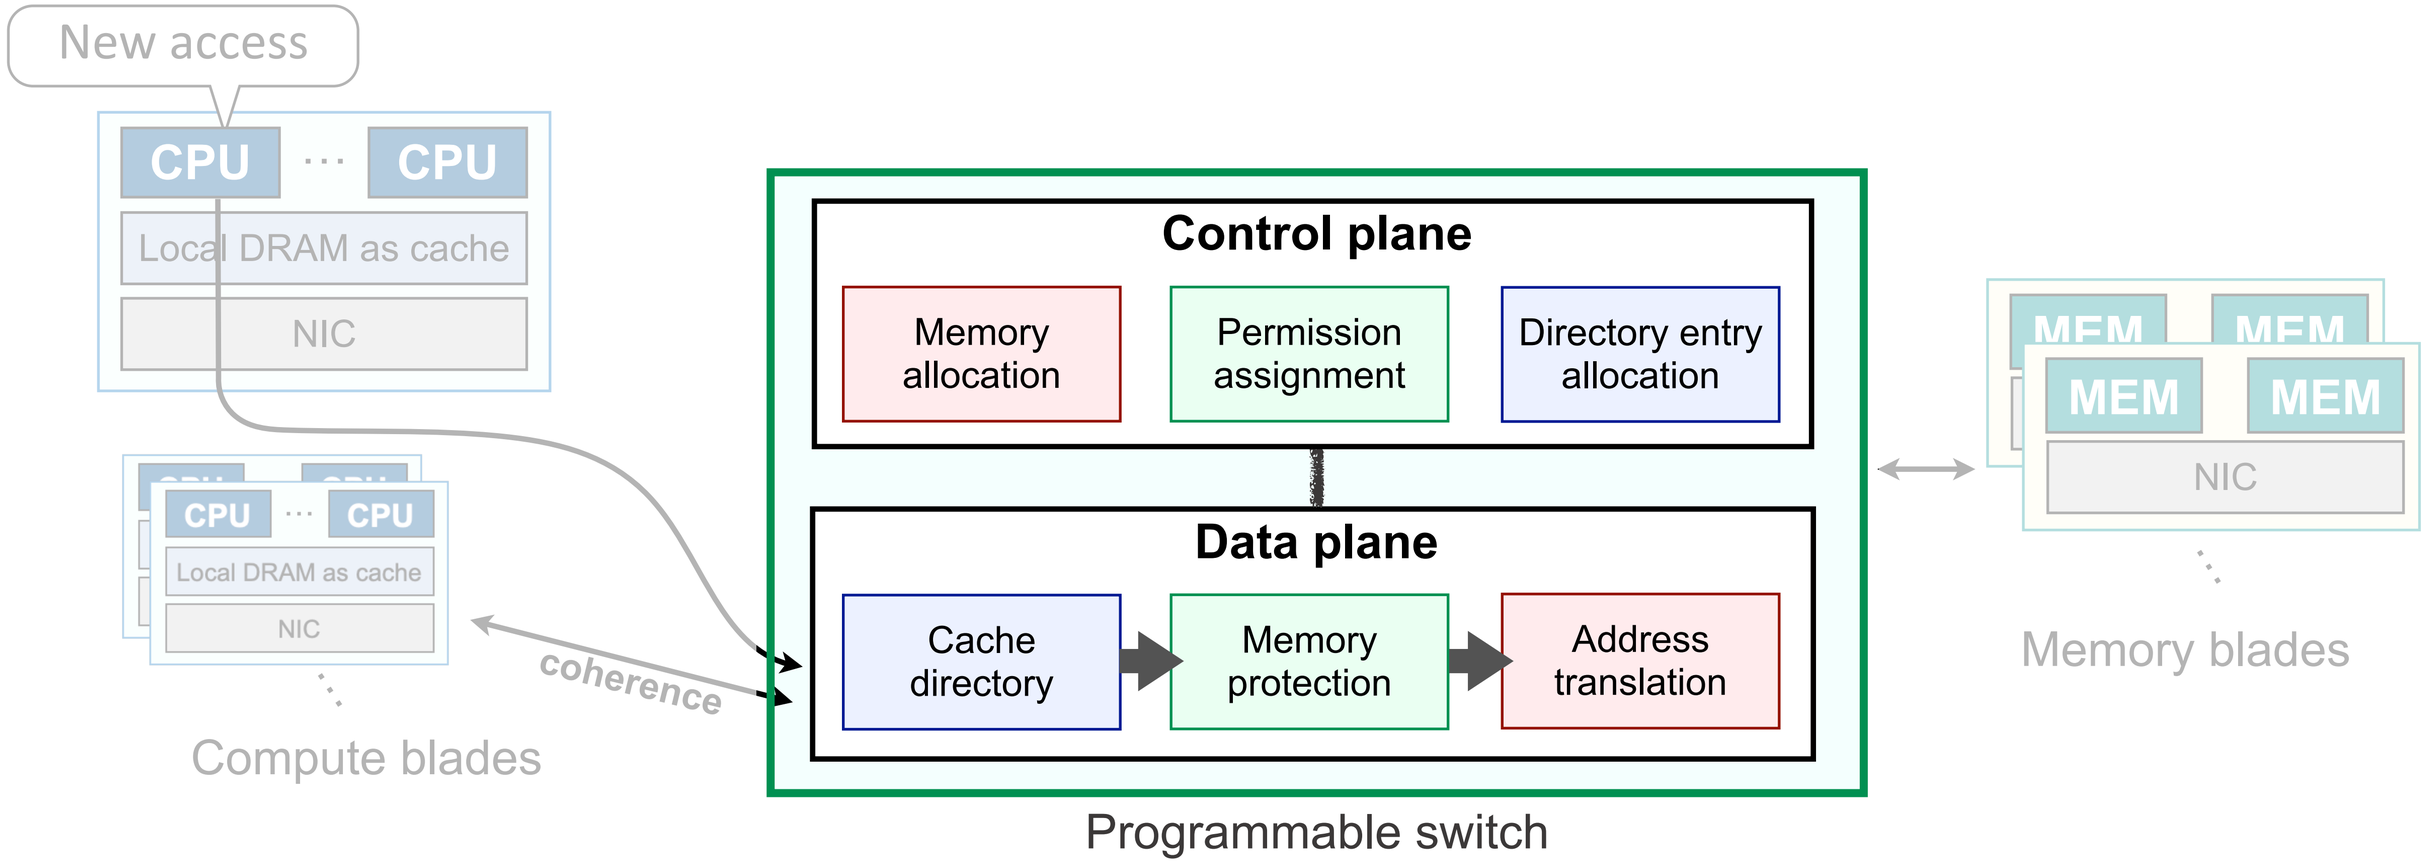

MIND overview: challenges and solutions

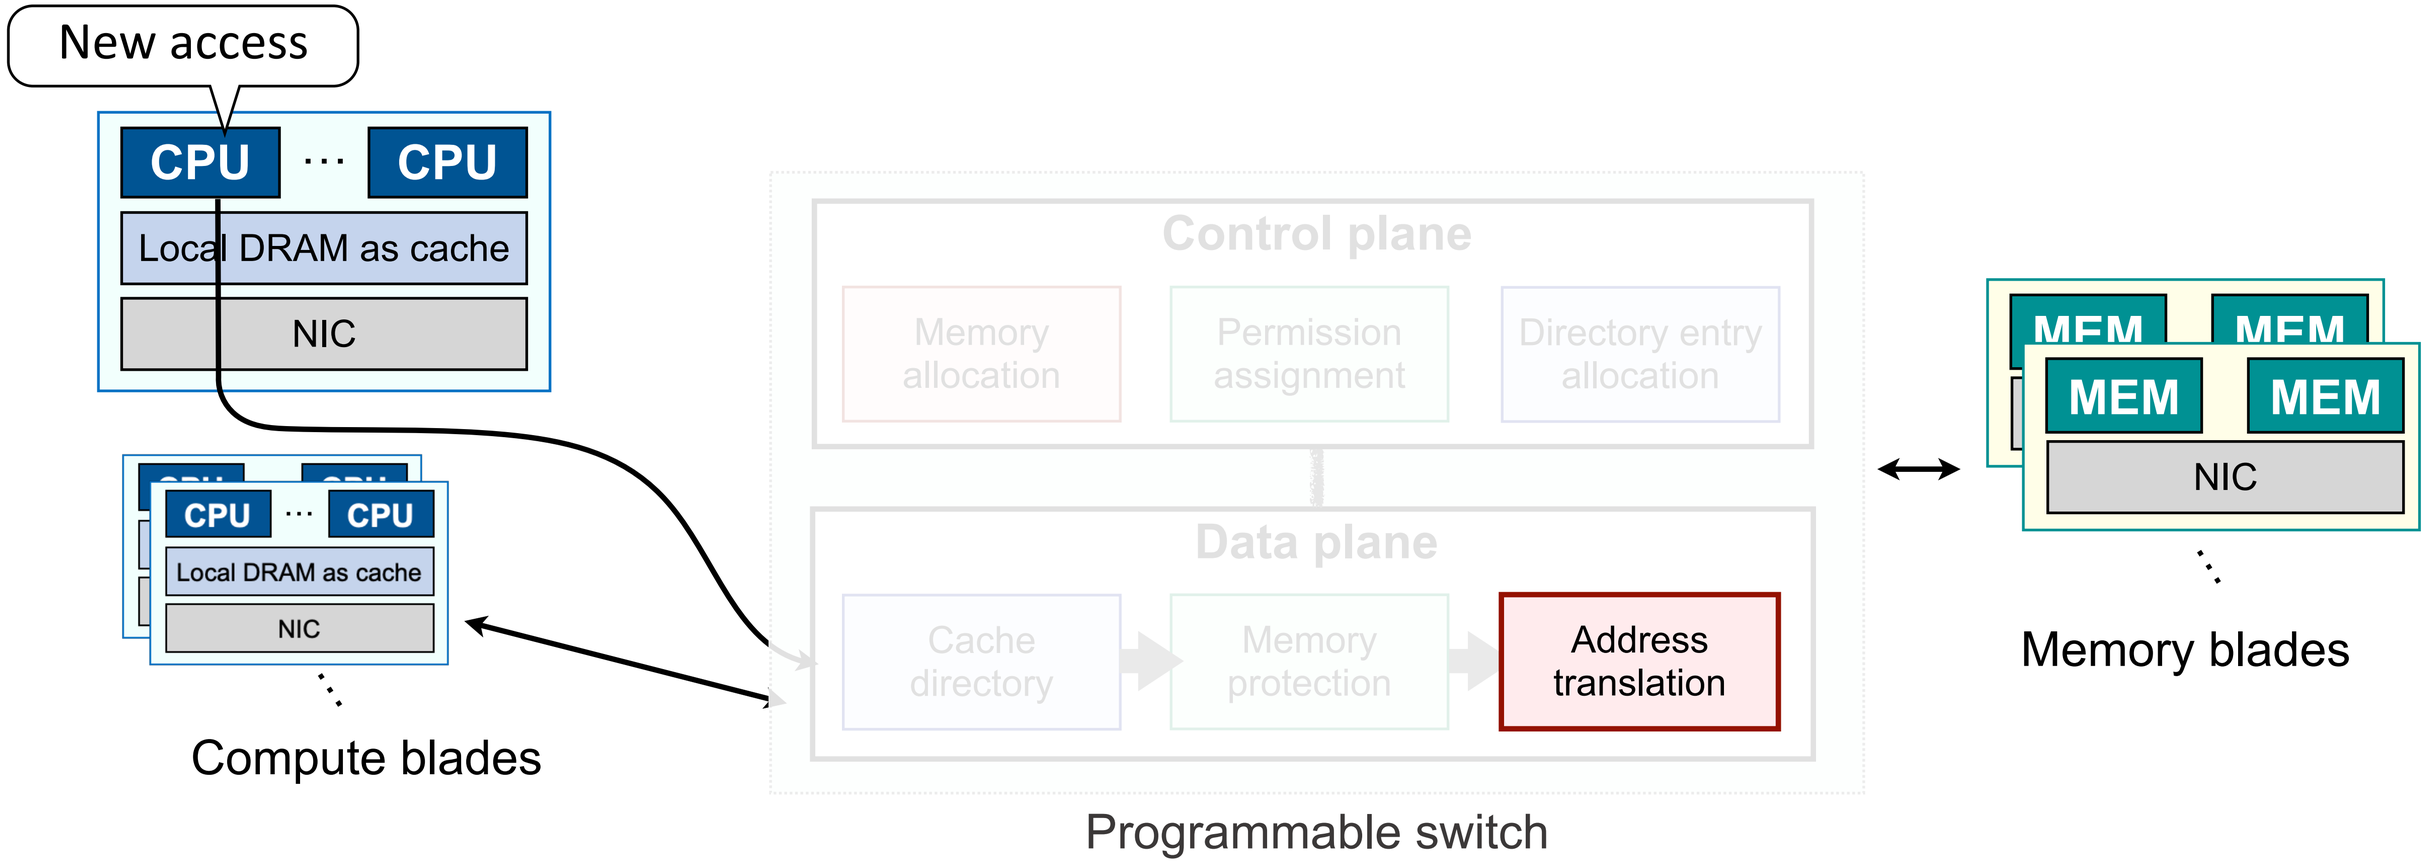

MIND overview: challenges and solutions

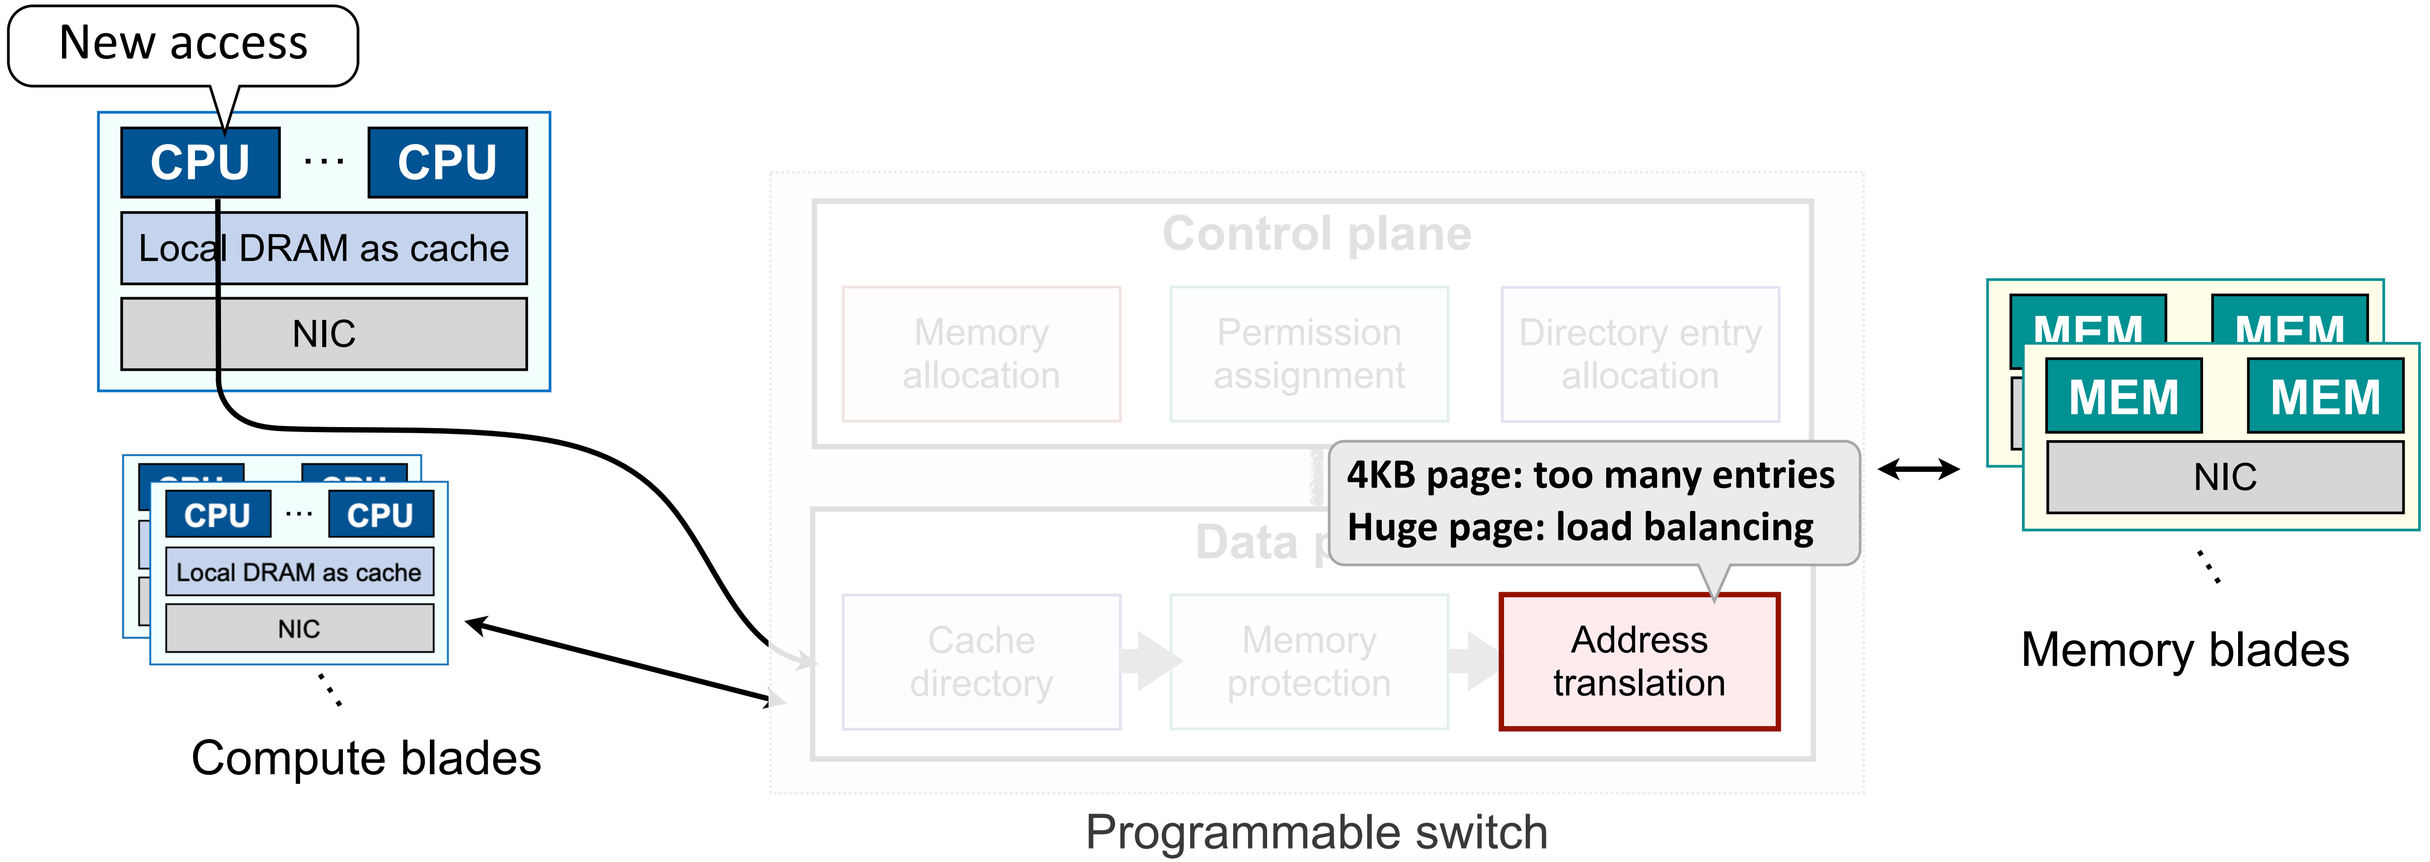

MIND overview: challenges and solutions

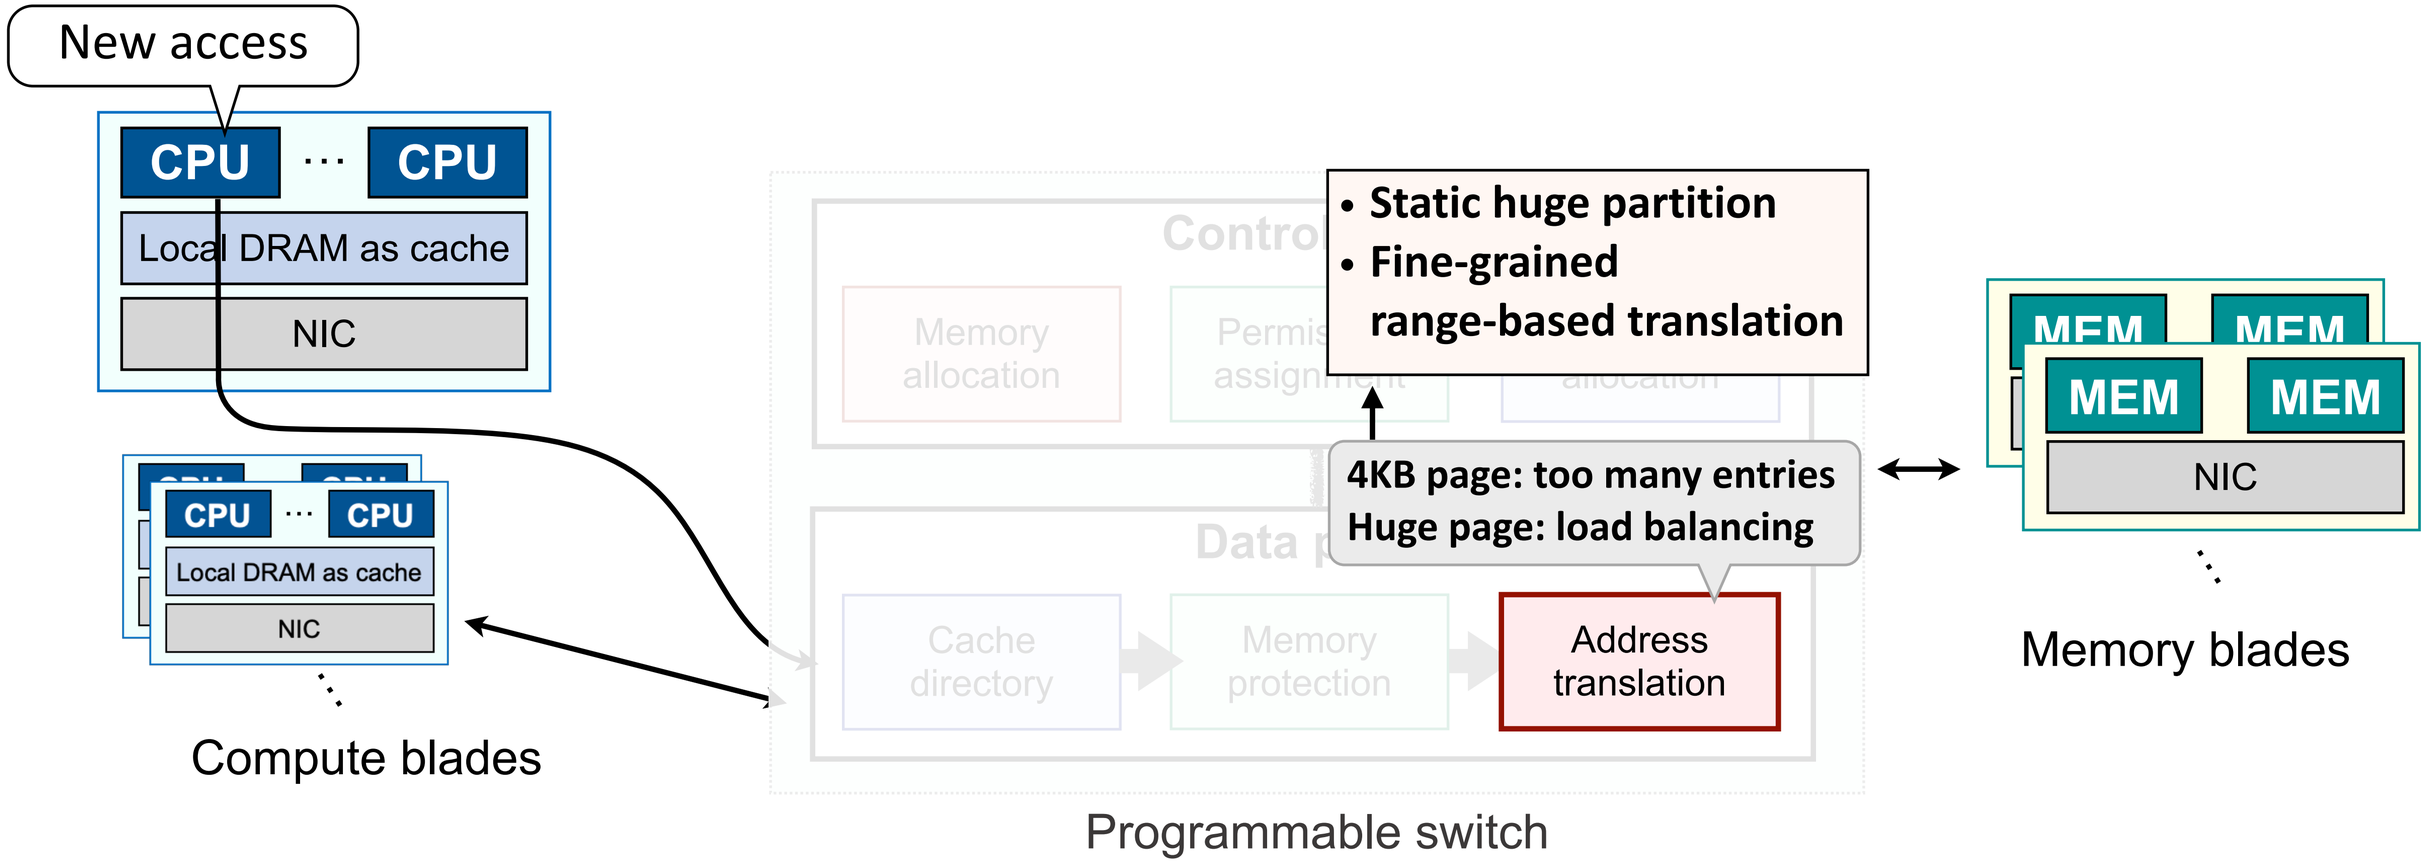

MIND overview: challenges and solutions

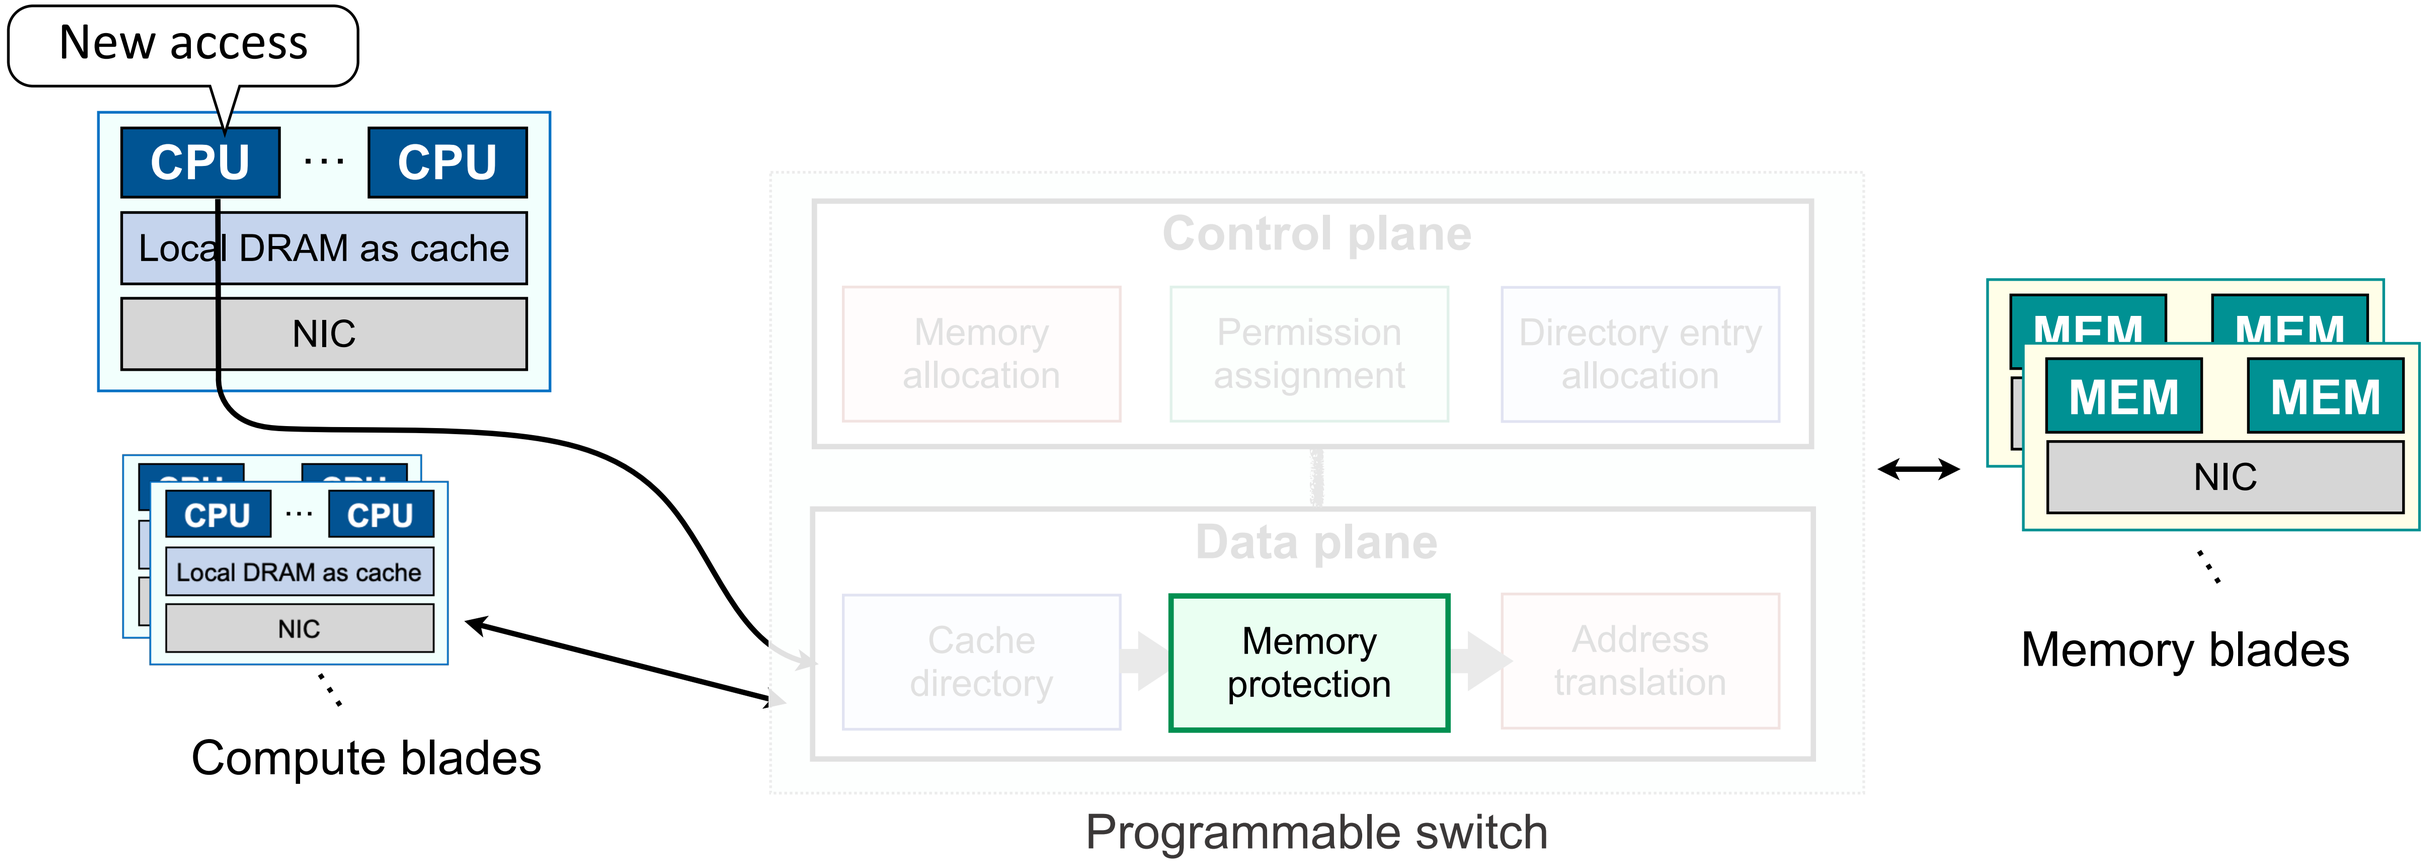

MIND overview: challenges and solutions

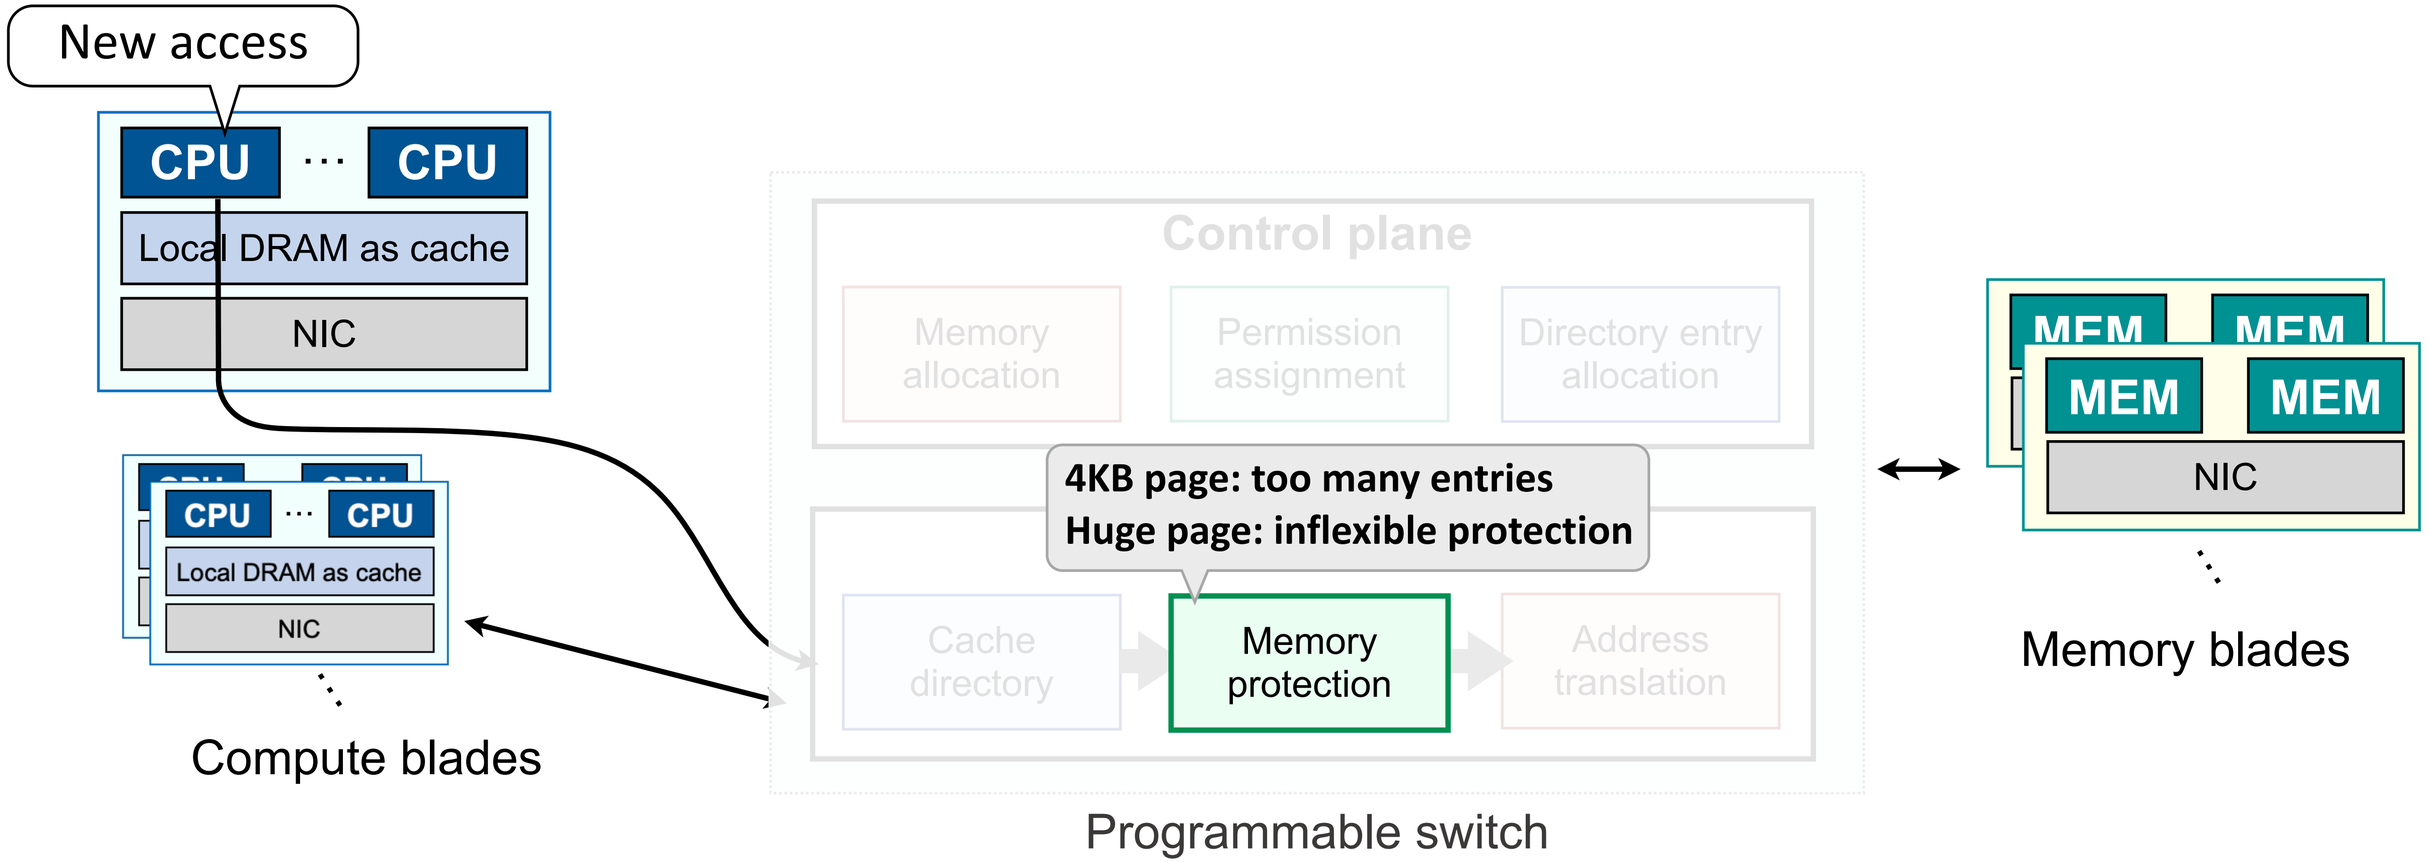

MIND overview: challenges and solutions

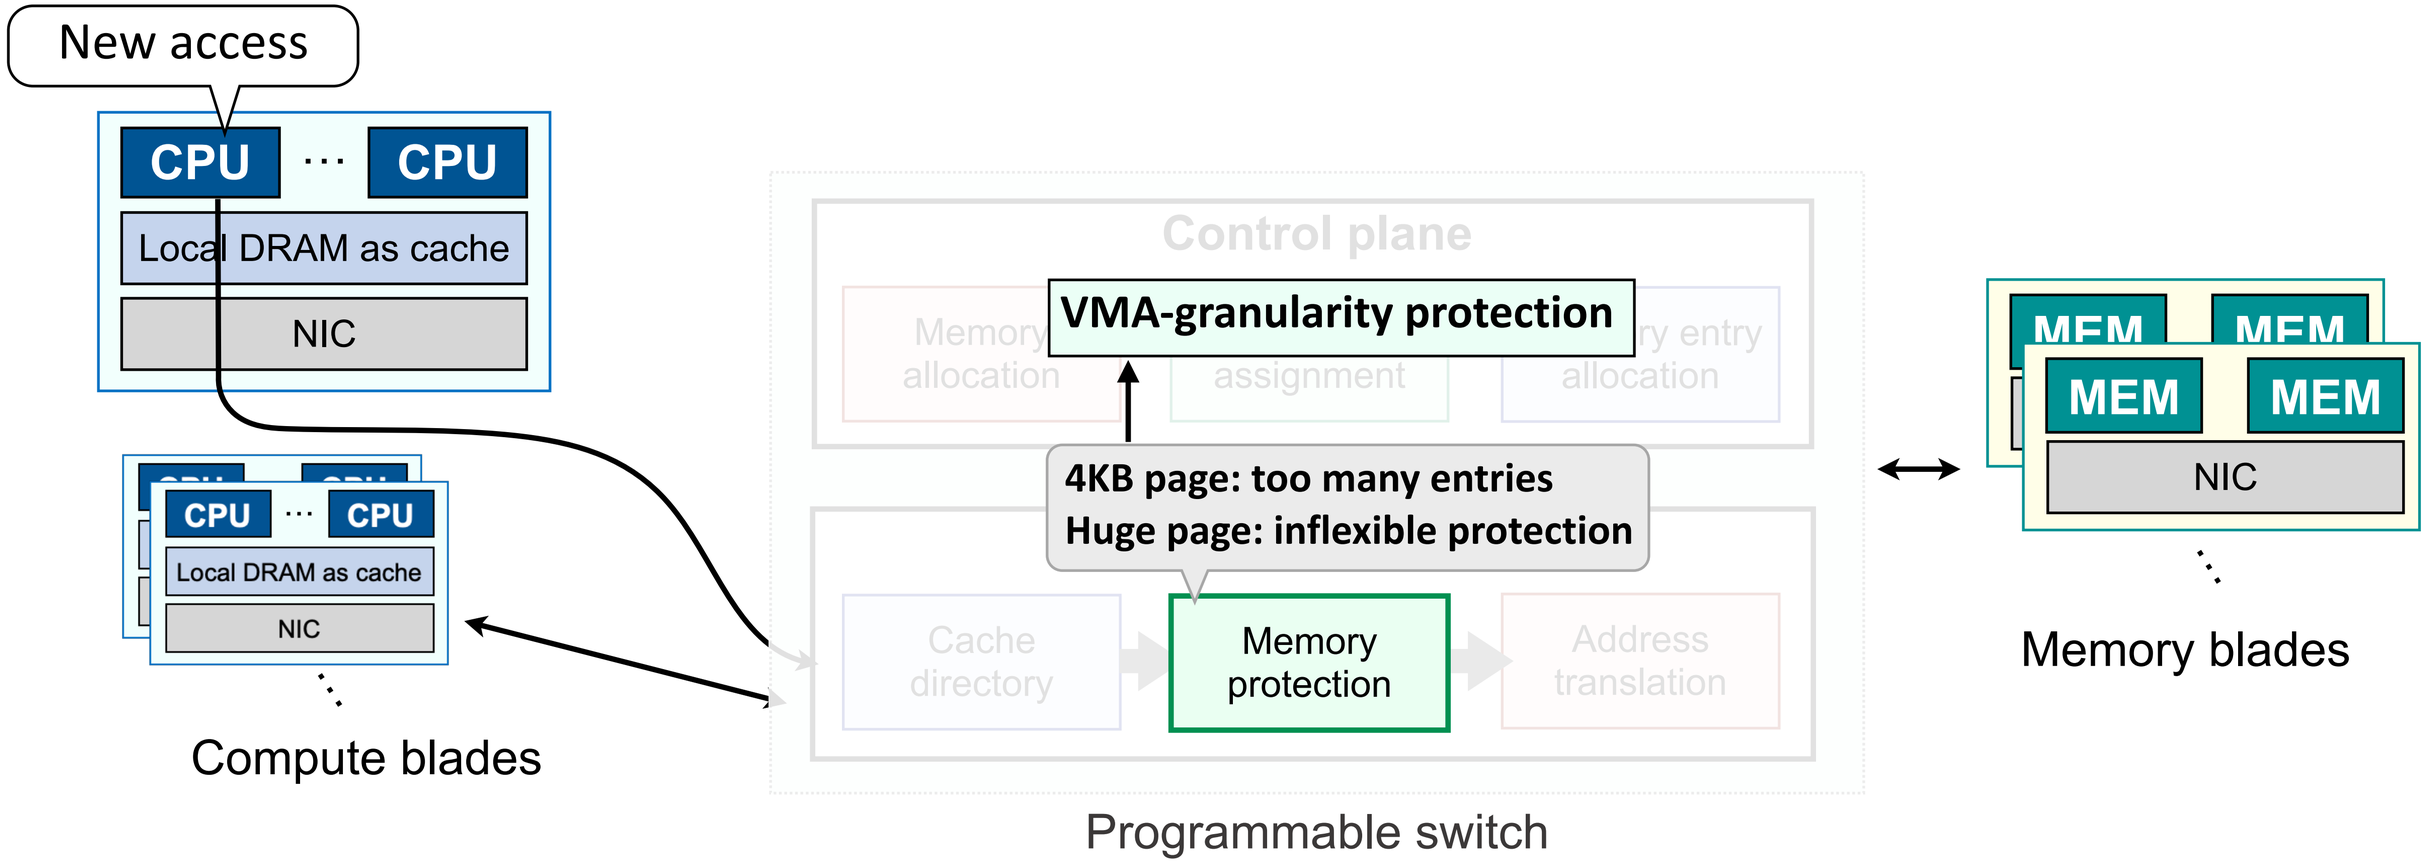

MIND overview: challenges and solutions

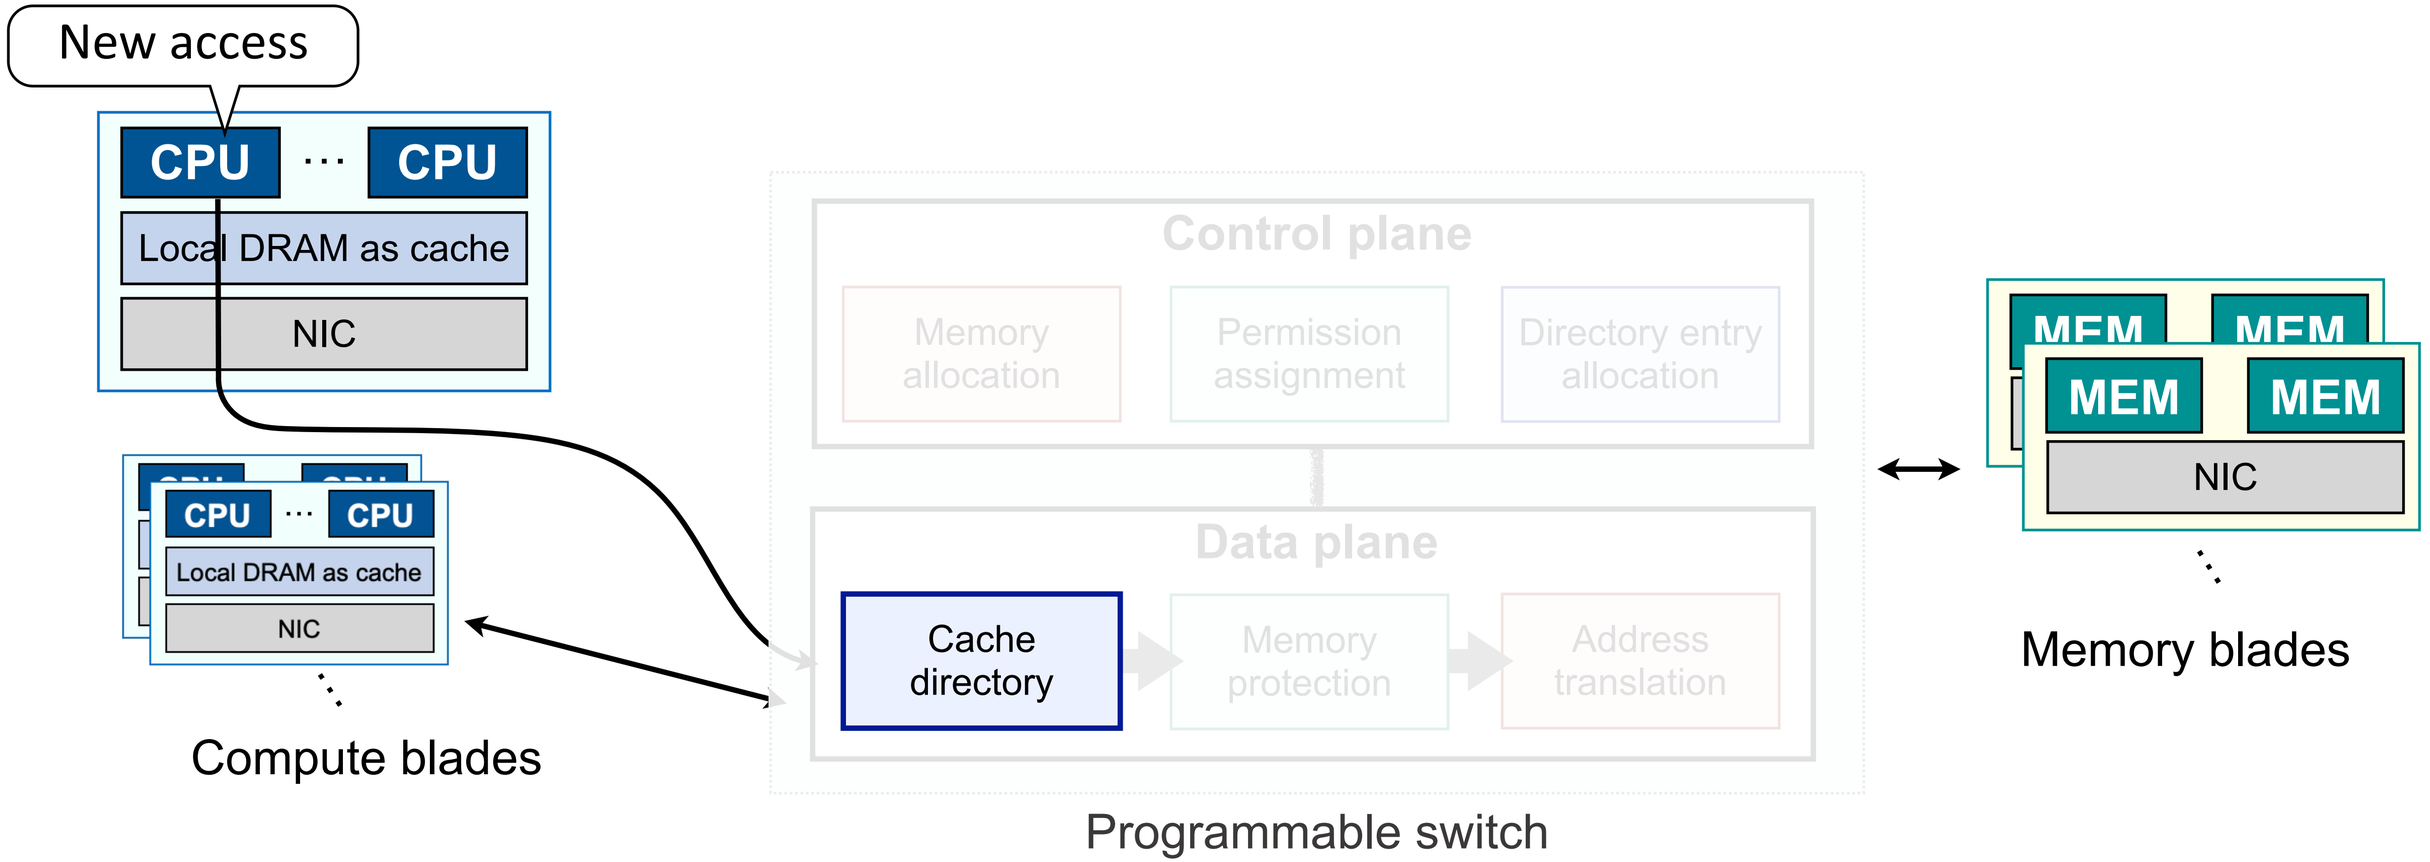

MIND overview: challenges and solutions

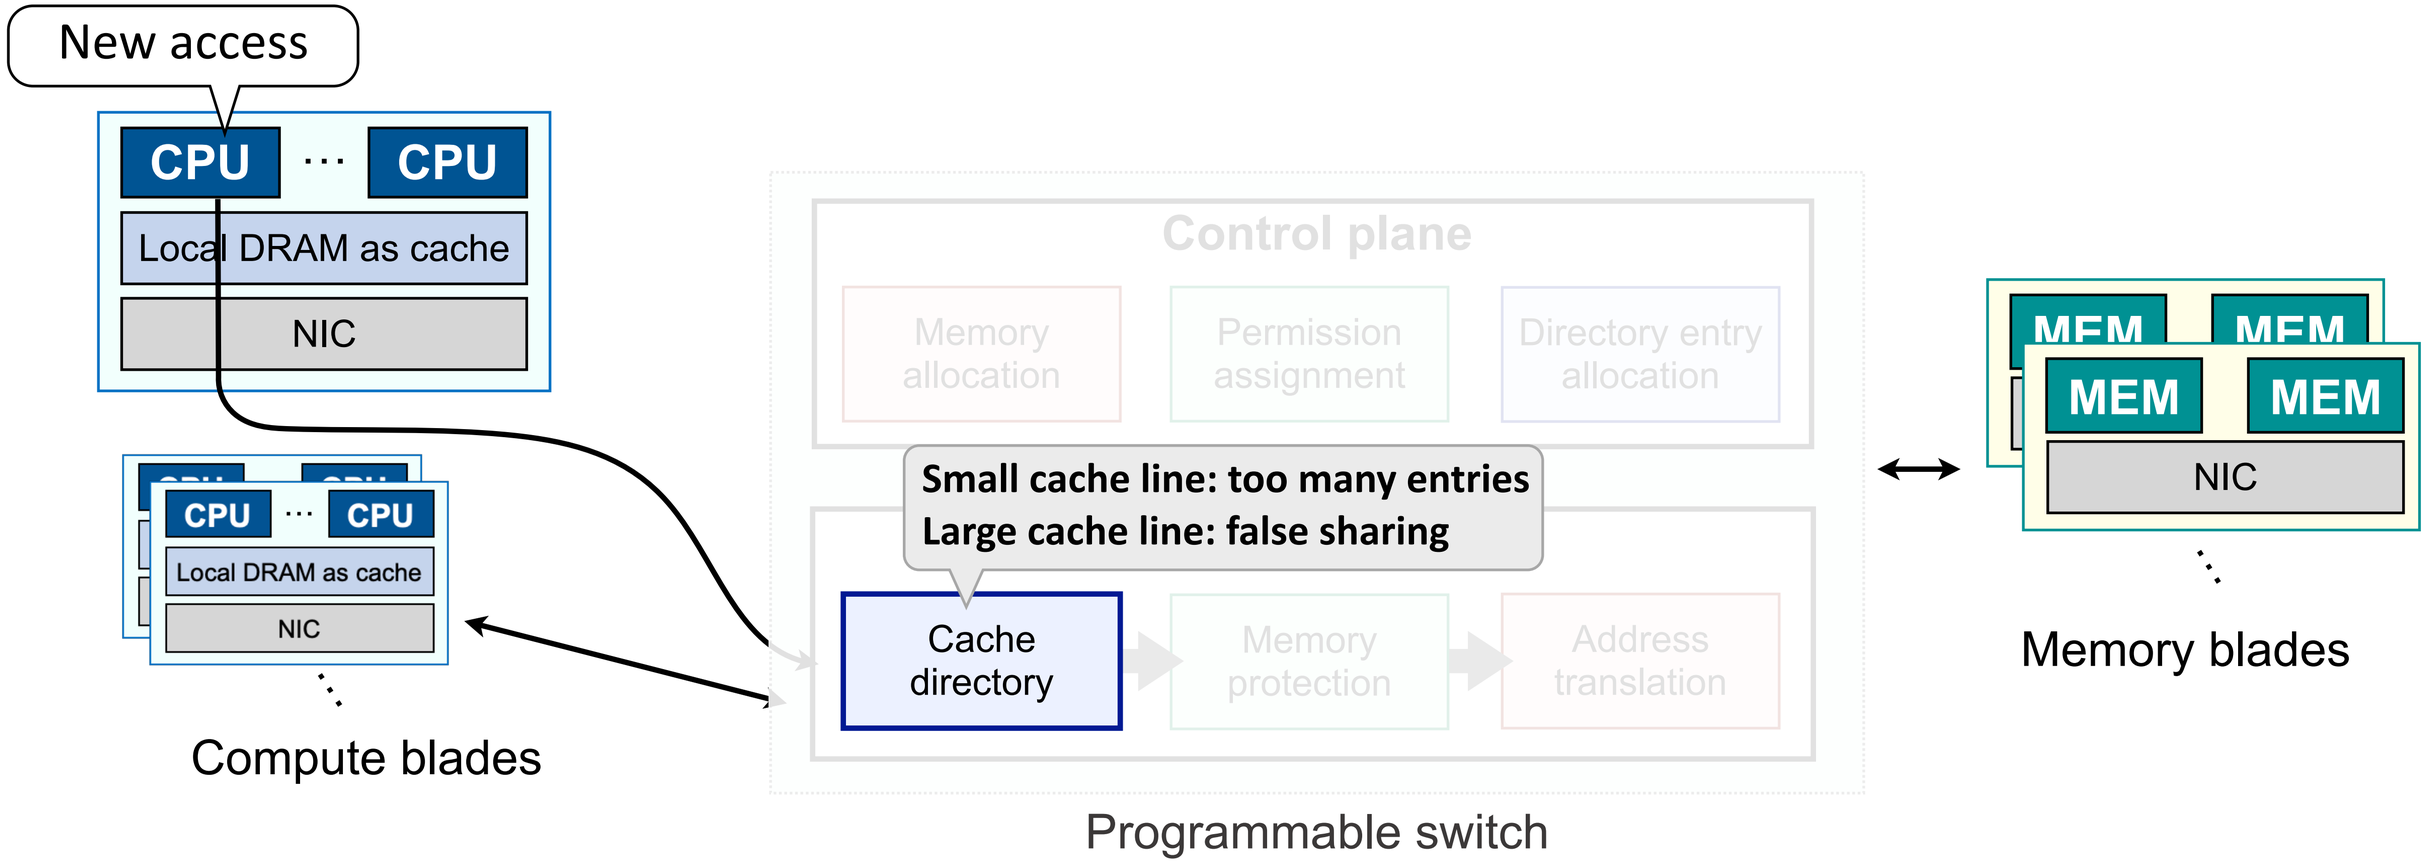

MIND overview: challenges and solutions

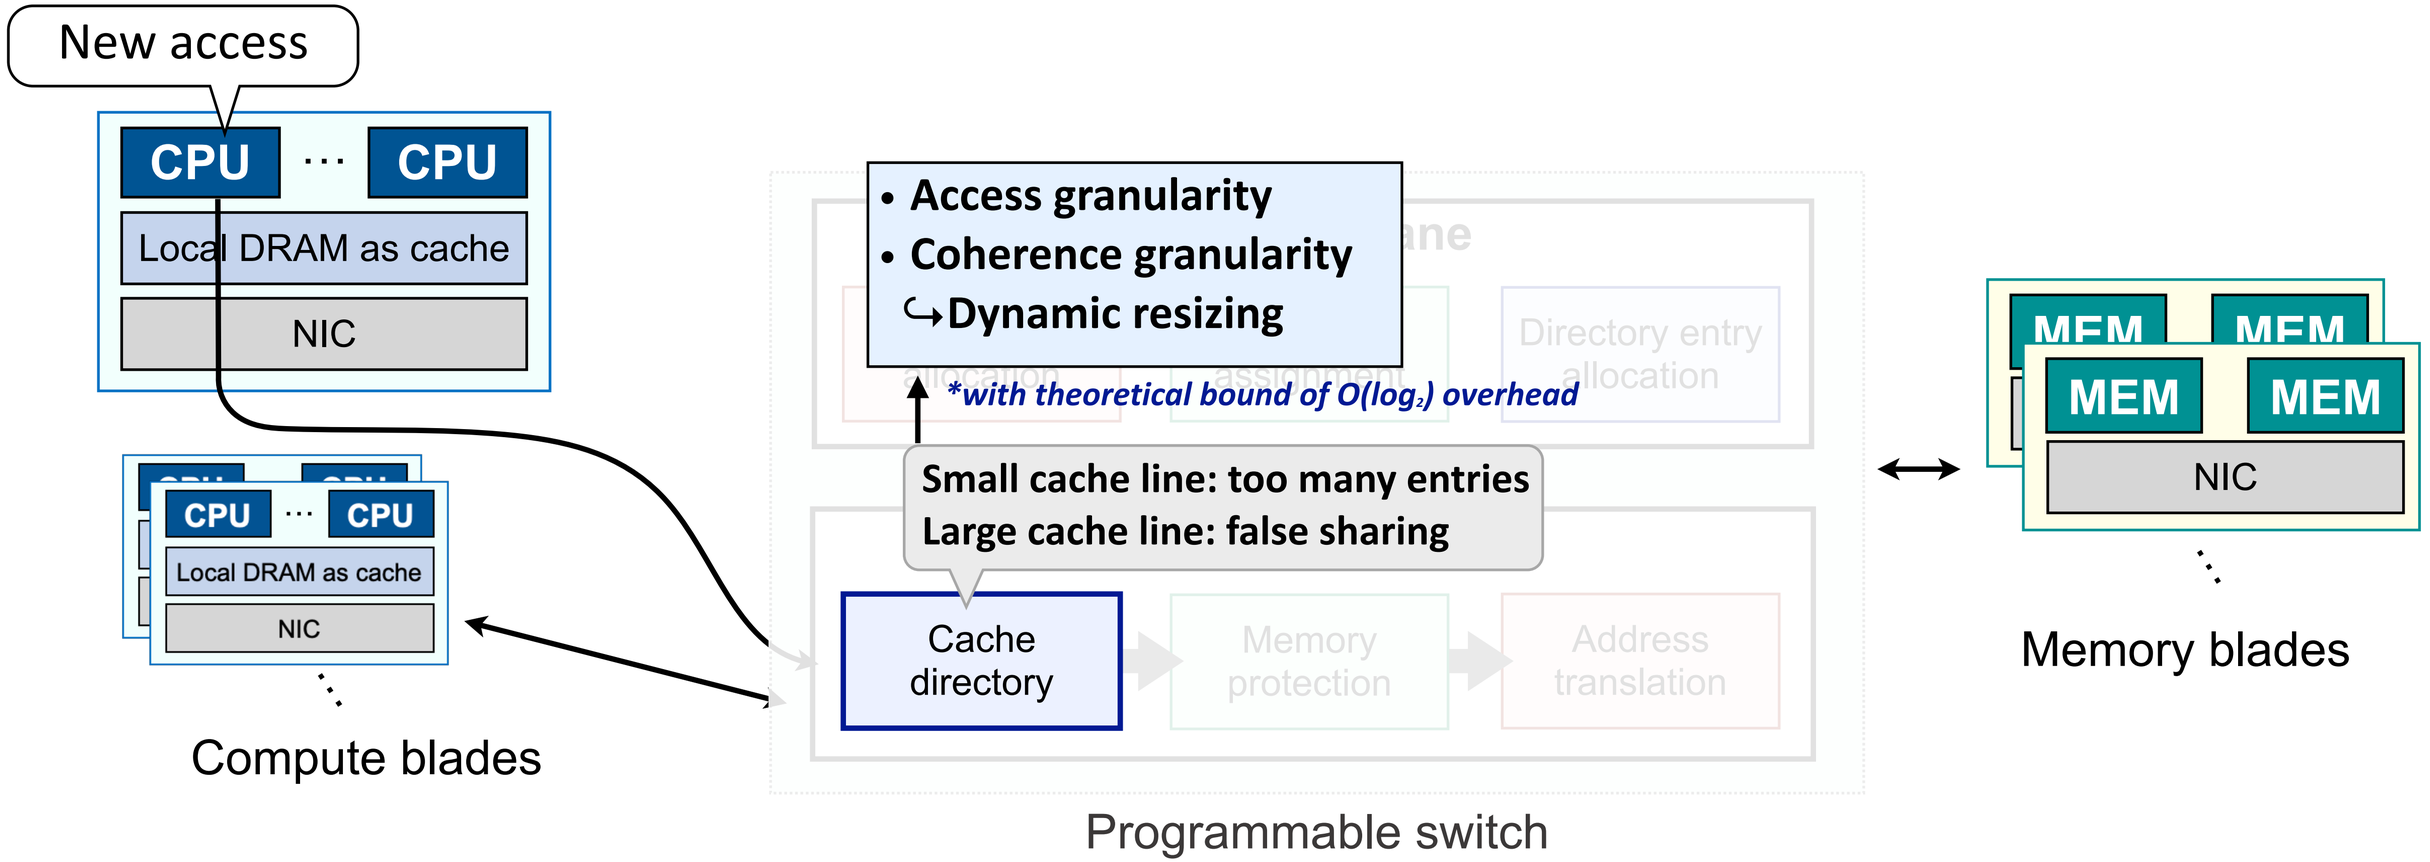

MIND overview: challenges and solutions

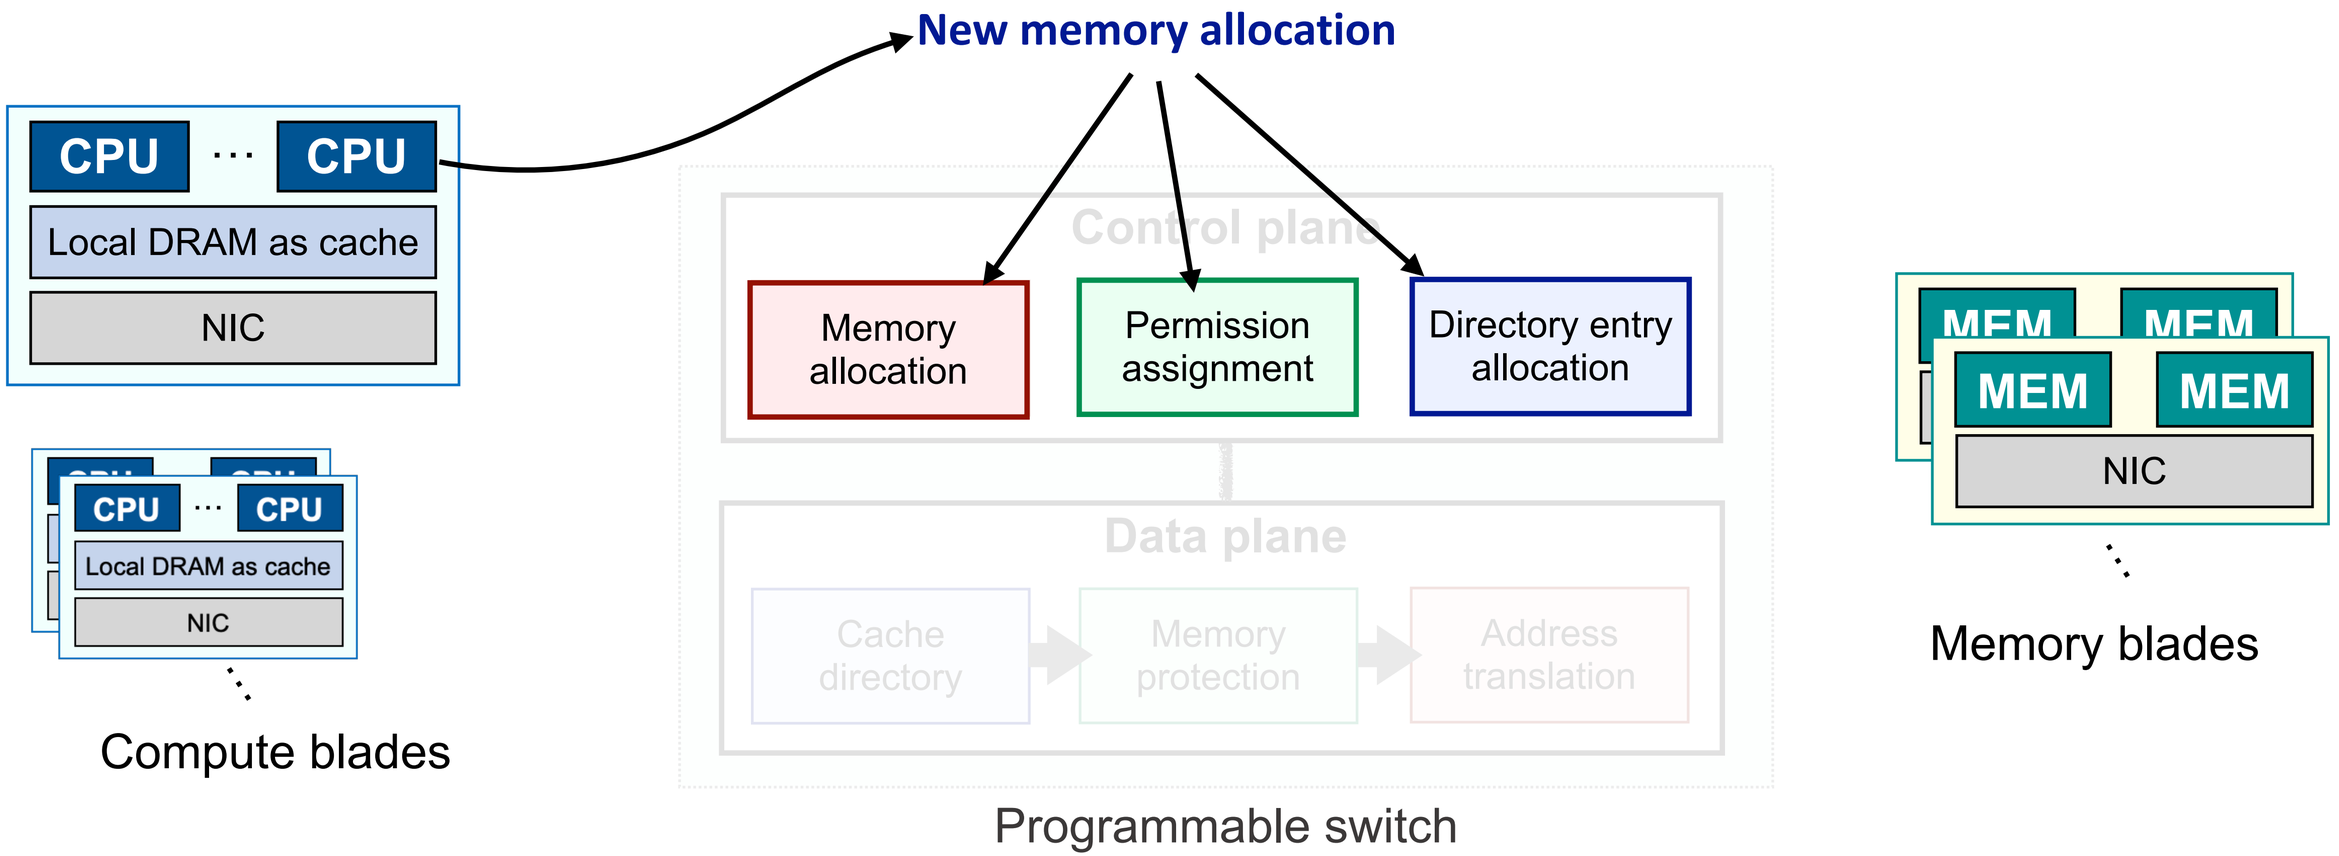

Performance Evaluation

- **Workloads with low contention**
  - TensorFlow (ResNet50), YCSB workload C (read only)

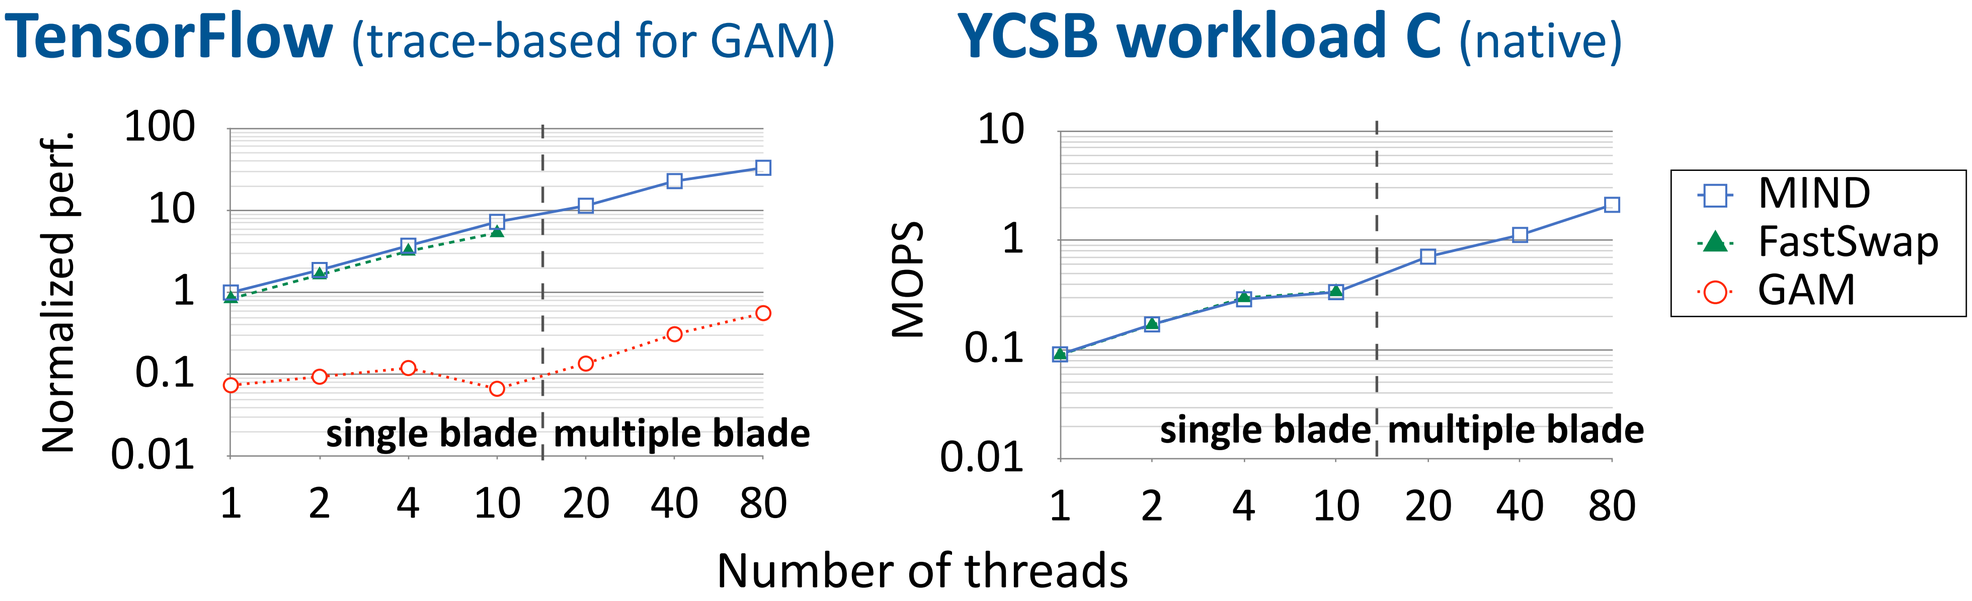

Performance Evaluation

- **Workloads with high contention**
  - GraphChi (PageRank), YCSB workload A (50 % update, 50 % read)

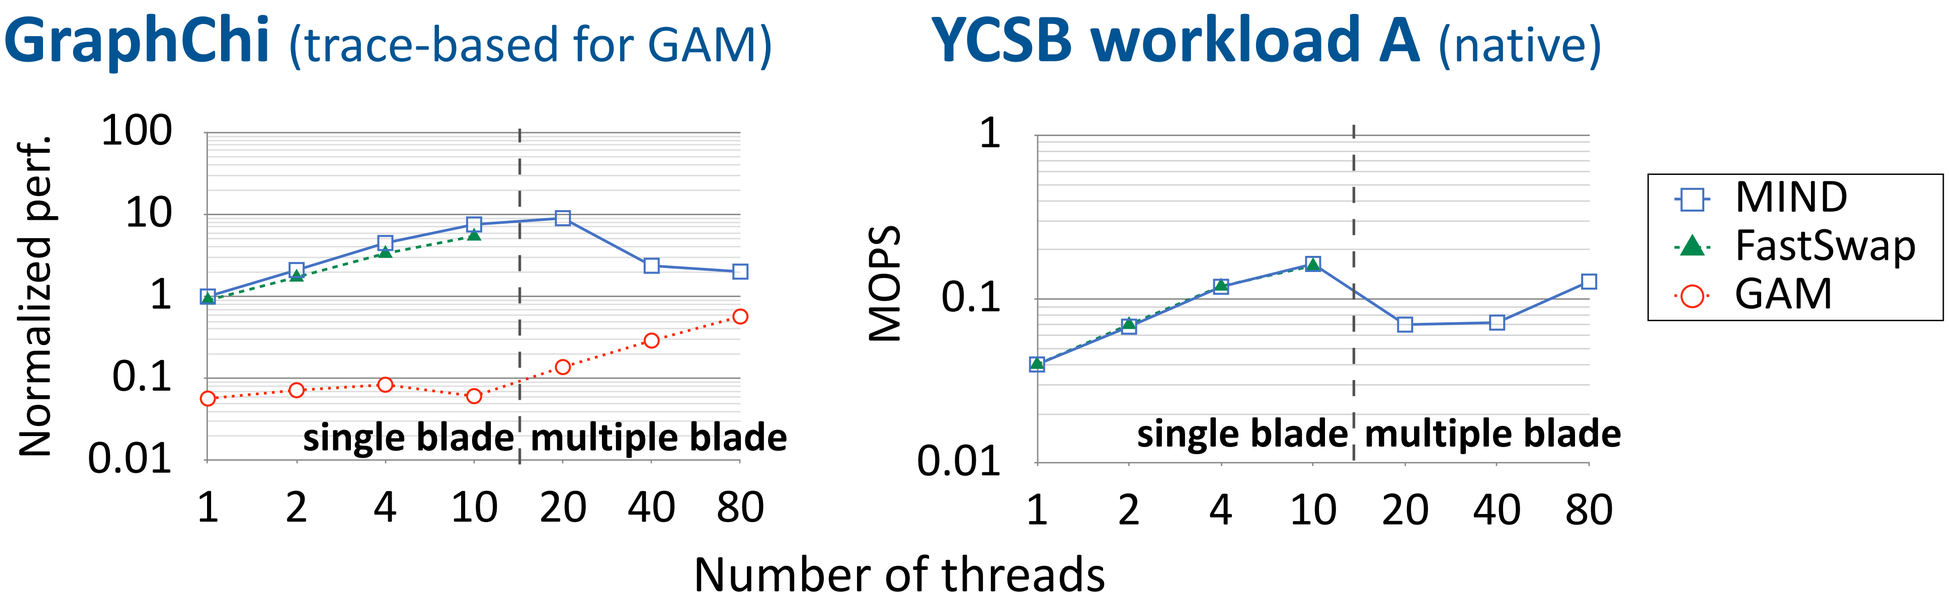

Detailed trade offs for design choices in the paper!

Conclusion & Summary

- **Trade-off between resource elasticity and performance** in memory disaggregation

Conclusion & Summary

- **Trade-off between resource elasticity and performance** in memory disaggregation

- We designed MIND, an **in-network MMU** by leveraging programmable network

Conclusion & Summary

- **Trade-off between resource elasticity and performance** in memory disaggregation

- We designed MIND, an **in-network MMU** by leveraging programmable network

- Our prototype of MIND can **match the performance of prior proposals and provide transparent elasticity**